

Converting the R Package `rpart` to Python: A Technical Report

Yufei Cai

2026

Contents

1	Introduction	4
1.1	Package Architecture	4
1.2	Target Package Layout	5
2	Phase 1: C Source Dependency Analysis	5
2.1	Motivation	5
2.2	Methodology	5
2.3	Key Findings	6
3	Phase 2: R External Item Extraction from C Sources	6
3.1	Motivation	6
3.2	Methodology	7
3.3	Key Findings	7
4	Phase 3: Fake C++ Header Implementation Guides	7
4.1	Motivation	7
4.2	Methodology	8
4.3	Classification of External Items	8
4.4	Key Technical Decisions	9
5	Phase 4: Fake C++ Header Generation and Compilation Bug Audit	10
5.1	Motivation	10
5.2	Header Generation	10
5.3	Post-Generation Bug Audit	11
6	Phase 5: Pure-C Entry-Point Generation and Fake-Header Migration	13
6.1	Motivation	13
6.2	Phase 5A: Entry-Point Generation	13
6.3	Phase 5B: Bulk Fake-Header Migration	14

7	Phase 6: R Package Structural Dependency Analysis	14
7.1	Motivation	14
7.2	Methodology	15
7.3	Key Findings	15
8	Phase 7: Python Build Infrastructure	15
8.1	Motivation	15
8.2	Build System Bugs and Fixes	16
8.3	CI/CD Pipeline	16
9	Phase 8: Language Dependency Analysis and Conversion Guide Generation	17
9.1	Phase 8.1: Language Dependency Extraction	17
9.1.1	Motivation	17
9.1.2	Methodology	17
9.1.3	CSV Encoding Defect	17
9.1.4	Key Findings	17
9.2	Phase 8.2: Conversion Guide Generation	18
9.2.1	Motivation	18
9.2.2	Methodology	18
9.2.3	Verification Edge Case	18
9.2.4	Representative Guide Complexity	18
10	Phase 9: R-to-Python Function Conversion and Package Assembly	19
10.1	Phase 9.1: Batch R-to-Python Function Conversion	19
10.1.1	Motivation	19
10.1.2	Methodology	19
10.1.3	Cross-Cutting Translation Decisions	19
10.1.4	Notable Per-File Decisions	20
10.2	Phase 9.2: Package Assembly, API Restructuring, and Test Validation	21
10.2.1	Motivation	21
10.2.2	Package Assembly	21
10.2.3	API Restructuring	21
10.2.4	Runtime Bugs in rpart.py	22
10.2.5	Build System Fix	22
10.2.6	Test Results	22
10.3	Phase 9.3: R Interpreter Item Bridge and cffi Migration	23
10.3.1	Motivation	23
10.3.2	Methodology	23
10.3.3	C Helper File	23
10.3.4	rpartcallback.py Rewrite	24
10.3.5	Results	24
A	Phase 1: C Dependency Analysis	24
A.1	Skill: /analyze-c-folder-dependencies	24
A.2	Agent: @analyze-c-file-dependencies	25

B	Phase 2: R External Item Extraction	27
B.1	Skill: /extract-c-folder-r-extern-items	27
B.2	Agent: @extract-c-file-r-extern-items	28
C	Phase 3: Fake Header Implementation Guides	29
C.1	Skill: /generate-r-extern-fake-guides	29
C.2	Agent: @generate-r-extern-fake-guide	31
D	Phase 4: Fake Header Generation	37
D.1	Skill: /generate-r-extern-fake-headers	37
D.2	Agent: @generate-r-extern-fake-header	39
E	Phase 5A: Pure-C Entry-Point Generation	43
E.1	Skill: /generate-r-extern-raw-entry-points	43
E.2	Agent: @generate-r-extern-raw-entry-point	45
F	Phase 5B: Fake-Header Migration	52
F.1	Skill: /modify-c-folder-with-fake-headers	52
F.2	Agent: @modify-c-file-with-fake-headers	56
G	Phase 6: R Package Structural Dependency Analysis	61
G.1	Skill: /analyze-r-folder-dependencies	61
G.2	Agent: @analyze-r-file-dependencies	62
H	Phase 8.1: Language Dependency Extraction	64
H.1	Skill: /extract-r-folder-language-dependencies	64
H.2	Agent: @extract-r-file-language-dependencies	65
I	Phase 8.2: Conversion Guide Generation	66
I.1	Skill: /generate-language-dependency-conversion-guides	66
I.2	Agent: @generate-language-dependency-conversion-guide	67
J	Phase 9.1: Batch R-to-Python Conversion	69
J.1	Skill: /convert-r-folder-to-python	69
J.2	Sub-skill: /convert-r-file-to-python	72
J.3	Agent: @convert-r-function-to-python	73
K	Phase 9.2: Package Assembly	75
K.1	Skill: /combine-python-functions-into-folder	75
K.2	Sub-skill: /combine-python-functions-into-file	78
L	Phase 9.3: R Interpreter Item Bridge	79
L.1	Skill: /add-r-interpreter-items	79
L.2	Agent: @add-r-interpreter-items	81

1 Introduction

The `rpart` package is one of the most widely used machine learning libraries in the R ecosystem. It implements the CART (Classification and Regression Trees) methodology introduced by Breiman, Friedman, Olshen, and Stone (1984), providing decision tree fitting for continuous, categorical, count, and survival responses through a unified interface. As a *recommended* package, `rpart` ships with every R installation and is depended upon by hundreds of CRAN packages. Despite the centrality of decision trees to modern machine learning, no faithful Python port of `rpart` exists that preserves its full feature set: surrogate splits for missing-data handling, four built-in splitting methods, a user-defined split callback mechanism, cost-complexity pruning, and the complete cross-validation and visualization infrastructure.

This project undertakes the systematic conversion of `rpart` version 4.1.27 to Python, targeting behavioral equivalence with the original R package. The target Python package is named `r2py_rpart`. Rather than reimplementing the algorithms from scratch—which would introduce independent bugs and divergences from the canonical implementation—the strategy is to reuse the original C source code directly and to convert the R source layer function by function in a structured, dependency-aware manner.

The conversion strategy has two parallel tracks. The first track concerns the compiled C back-end: `rpart`’s five C entry points (`C_rpart`, `C_pred_rpart`, `C_xpred`, `C_rpartexp2`, `C_init_rpcallback`) are wrapped with a layer of fake C++ headers that replace the R C API (normally provided by `libR.so`) with standalone implementations, enabling the original C code to be compiled as a shared library and called from Python via `cffi`. The second track concerns the R source layer: all 47 functions across 36 R source files are analyzed for their dependency structure and converted to Python one by one in a topologically sorted order, with each conversion guided by pre-generated Markdown guides for the 172 distinct base-R language constructs used in the package.

The entire conversion process is divided into sequential phases. Phases 1–3 analyze the existing codebase. Phases 4–6 build the C compilation infrastructure. Phase 7 analyzes the R source layer. Phase 8 establishes the Python packaging infrastructure. Phases 9–10 drive the R-to-Python function conversion and package assembly. This report documents each phase in order, describing both what was accomplished and the technical reasoning behind the design choices made.

1.1 Package Architecture

`rpart` follows a three-tier architecture. The public R API (Tier 1) provides formula-based model fitting, S3 method dispatch, and diagnostic utilities. The R infrastructure layer (Tier 2) handles data preparation, method-specific initialisation, coordinate geometry for tree plotting, and numeric formatting. The C back-end (Tier 3) performs the computationally intensive work: greedy recursive partitioning, surrogate-split search, cross-validation across the full complexity-parameter sequence, and prediction routing. The C code is organized around a dispatch table (`func_table.h`) that maps method integers to four function-pointers

slots (`init_split`, `choose_split`, `eval`, `error`), keeping all bookkeeping logic completely method-agnostic. Method integer 4 is reserved for user-defined (R-level callback) splits, implemented via the `rpartcallback` mechanism.

A key structural observation motivating the conversion strategy is that the R-to-C interface is narrow: only five `.Call()` invocations exist across all 36 R source files. This means that if these five entry points can be faithfully exposed through a Python-callable interface, the R source layer can be converted independently without touching the C layer.

1.2 Target Package Layout

The `r2py_rpart` Python package is organized as follows. The directory `r2py_rpart/src/` contains a copy of the original C source files from `rpart/src/`, modified to replace real R API includes with fake headers. The directory `r2py_rpart/r_fake_headers/` contains the 53 fake C++ header files. The directory `r2py_rpart/c_entry_points/` contains the six pure-C wrapper files that translate between the Python-callable (plain-array) interface and the SEXP-based internal interface. The directory `r2py_rpart/r2py_rpart/` contains the Python package itself: 37 `.py` modules converted from the R source, plus `__init__.py` which loads `_rpart_core.so` via `cffi` and exposes the low-level C entry points as private symbols.

2 Phase 1: C Source Dependency Analysis

2.1 Motivation

Before any code can be converted or the C compilation infrastructure can be designed, it is necessary to understand the internal structure of the C source layer: which functions call which other functions, and which functions call into the R C API. This structural knowledge serves two purposes. First, it provides the call graph needed to determine the correct bottom-up order for conversion (leaf functions with no internal dependencies can be translated first, followed by their callers). Second, it identifies exactly which R API symbols are used, which is the foundation for the fake-header generation in later phases.

2.2 Methodology

The `/analyze-c-folder-dependencies` (see Appendix A.1) skill was invoked against the `rpart/src/` directory, which contains all 35 C and header files of the package:

```
/analyze-c-folder-dependencies rpart/src/
```

This skill dispatched the `@analyze-c-file-dependencies` sub-agent sequentially to each file. Each agent read the full source text of its target file and all transitively included local headers, then classified every function call as either an *internal dependency* (resolved within the project) or an *external dependency* (R API call such as `R_alloc`, `PROTECT`, `error`). The output was a JSON file per C source, saved to `c_refactor_analysis/src/` with the naming convention `<filename>.json`.

Following the automated analysis, a manual cross-examination of all 35 JSON files revealed three inconsistencies that were corrected:

1. `xpred.c.json`: the function pointers `rp_init` and `rp_eval` were missing from `xpred`'s internal dependencies despite being called via `(*rp_init)(...)` and `(*rp_eval)(...)`.
2. `rpart_callback.c.json`: the `_` macro was incorrectly listed as a direct internal dependency of `init_rpcallback`; it is only called inside `compat_getVar`, not `init_rpcallback` itself.
3. `rpart_callback.c.json`: the `_` macro entry lacked its `dgettext` external dependency, which was correct in the analogous entry in `rpart.h.json`.

Running the graph-construction script `dependency_graph.py` after corrections revealed five identifiers referenced as internal dependencies but absent as defined entries in any file: the four `EXTERN` function pointer variables `rp_init`, `rp_choose`, `rp_eval`, `rp_error` (declared in `rpart.h`) and the `static` function pointer `impurity` (in `gini.c`). These are runtime-dispatch function pointer variables, not function definitions; they were added with empty dependency lists to resolve the dangling graph nodes.

The `networkx` package (version 3.6.1) was installed into the `r-to-python` conda environment as a missing dependency before the final script runs.

2.3 Key Findings

The final dependency graph contains 136 nodes (35 file nodes, 66 function/macro nodes, 35 external-dependency nodes) and 255 edges. The topological level distribution in `dependency_levels.csv` shows 38 nodes at level 0 (entry points or leaves with no unresolved callers), 24 at level 1, and 4 at level 2. Of the 66 function nodes, 47 are pure leaves that make no internal calls. Structurally, `rpart.c` is the primary integration point (12 internal dependencies, 15 R API calls), while `rpart_callback.c` is architecturally isolated—it does not include `rpart.h` and implements its own version-compatibility logic for $R < 4.5.0$ vs. $R \geq 4.5.0$ using its own `compat_getVar` shim.

3 Phase 2: R External Item Extraction from C Sources

3.1 Motivation

Having established the internal call graph in Phase 1, the next prerequisite for the C compilation strategy is a precise inventory of every R API symbol used by the package: which macros, types, and functions from R's header files appear in the C source, in which files, at which line numbers, and in what syntactic context. This inventory serves as the input to the fake-header generation pipeline: each unique R external item requires its own fake implementation. Without a complete and accurate inventory, fake headers would be incomplete and the compilation would fail at unpredictable locations.

3.2 Methodology

The `/extract-c-folder-r-extern-items` (see Appendix B.1) skill was invoked against `rpart/src/`:

```
/extract-c-folder-r-extern-items rpart/src/
```

The `@extract-c-file-r-extern-items` sub-agent was dispatched for each of the 35 C and header files. Each agent read the source file and all transitively included local headers, then searched the R include directory at `~/conda/envs/r-to-python/lib/R/include/` to determine whether each non-standard identifier was declared in an R header. A critical disambiguation was applied throughout: local wrapper macros defined in `rpart.h` itself (`ALLOC`, `CALLOC`, `RPARTNA`, `LEFT`, `RIGHT`) were consistently excluded, since they are declared in the project header, not in the R include directory, even though they wrap R API calls (`R_alloc`, `R_chk_calloc`, `ISNAN`).

Per-file CSVs were saved to `r_extern_analysis/src/` with the schema `external_item`, `header_file`, `category`, `file_name`, `line_number`, `context_statement`. A combination script `r_extern_analysis/combine_csvs.py` merged all files into a single sorted master table `r_extern_analysis/combined.csv`.

3.3 Key Findings

The combined table contains 235 reference rows across 51 unique R external identifiers drawn from 16 source files (the remaining 19 files rely exclusively on project-internal headers and standard C primitives). The 51 identifiers are drawn from 10 R include headers, with `Rinternals.h` accounting for 165 of the 235 references. The category breakdown is 184 functions (78%), 39 types (17%), and 12 variables (5%).

Two files stand out architecturally. `rpart.c` and `xpred.c` are the most R-API-dense, making heavy use of `SEXP`, `PROTECT/UNPROTECT`, `allocVector/allocMatrix`, `INTEGER`, `REAL`, and list-construction functions. `rpart_callback.c` is uniquely notable: it is the only file that calls R interpreter-level functions (`eval`, `findVar`, `findVarInFrame`, `install`, `R_getVar`). These five symbols are required only when `rpart()` is called with a user-defined splitting method (`method=4`); all four standard methods (`anova`, `poisson`, `class`, `exp`) execute without them. This observation will become the basis for the Category E design in Phase 3.

4 Phase 3: Fake C++ Header Implementation Guides

4.1 Motivation

The central challenge in making the `rpart` C source files callable from Python without `libR.so` is that they include R's C API headers and call R's C API functions. These headers and functions are not available without a running R interpreter. The solution adopted in this project is to write *fake* C++ headers that provide compilable, self-contained implementations of all 51 R external items identified in Phase 2. When the C source files include `fake_R.h` instead of the real R headers, they can be compiled and linked as a standalone shared library.

Before writing the headers themselves, it is essential to plan the implementation strategy for each item. Memory management, type representations, and singleton semantics all require careful design to ensure that the 36 source files compile together without ODR (One Definition Rule) violations or linkage conflicts. Phase 3 produces one Markdown guide per external item, documenting the full implementation blueprint.

4.2 Methodology

The `/generate-r-extern-fake-guides` (see Appendix C.1) skill was invoked with the master CSV as input:

```
/generate-r-extern-fake-guides rpart/src/ r_extern_analysis/combined.csv  
r_extern_analysis/fake_guides/
```

The `@generate-r-extern-fake-guide` sub-agent was invoked sequentially for each of the 51 unique external items, in the order they appear in the pre-sorted CSV. Sequential processing is mandatory here because foundational definitions (the `SEXP` struct layout, the `SEXP_TYPE` enum, the arena allocator) must be established before dependent items (`INTEGER`, `allocVector`, `error`) can reference them. Each agent read the relevant `rpart` source files and the real R include headers, and produced a Markdown guide named `<external_item>.md` in `r_extern_analysis/fake_guides/`.

4.3 Classification of External Items

The 51 guides were classified into five categories based on the fake implementation strategy required:

Category	Description	Items
A (19)	Type, enum constant, or registration no-op	SEXP, INTSXP, REALSXP, STRSXP, VECSXP, DL_FUNC, DllInfo, R_CallMethodDef, Rboolean, TRUE, FALSE, R_Version, R_NilVal, R_NamesSymbol, R_UnboundValue, R_forceSymbols, R_registerRoutine, R_useDynamicSymbols
B (16)	Accessor macro or inline function	INTEGER, REAL, CHAR, PROTECT, UNPROTECT, LENGTH, PRINTNAME, ISNAN, R_FINITE, asInteger, asReal, isReal, ncols, nrows, SET_STRING_ELT, SET_VECTOR_ELT
C (7)	Allocation or memory management	allocVector, allocMatrix, R_alloc, R_chk_calloc, R_Free, mkChar, setAttrib
D (4)	Error, warning, or print function	error, warning, Rprintr, R_CheckUserInterrupt
E (5)	R interpreter item (Python function-pointer bridge)	eval, findVar, findVarInFrame, install, R_getVar

4.4 Key Technical Decisions

Memory model. The guides establish a two-tier memory model. Scratch allocations via `R_alloc`/`ALLOC` delegate to a per-frame arena allocator (`ArenaFrame`) that is freed automatically when the `.Call` boundary wrapper returns. Persistent allocations for `SEXP` nodes and their data buffers live on the process heap and are freed explicitly via `R_Free` or `free_sexp()`.

R version pinning. `R_VERSION` must be faked as `R_Version(4,4,0) = 263168`. This value simultaneously satisfies two conflicting preprocessor guards: `R_VERSION >= R_Version(2,16,0)` in `init.c` (enabling the `R_forceSymbols` branch) and `R_VERSION < R_Version(4,5,0)` in `rpart_callback.c` (selecting the `compat_getVar` shim instead of the native `R_getVar` symbol, which would require `libR.so`).

Category E items. The five R interpreter items (`eval`, `findVar`, `findVarInFrame`, `install`, `R_getVar`) are only reachable when `rpart()` is called with `method=4` (user-defined splits). For the `install` function, the guides specify a self-contained C++ implementation using a `thread_local std::unordered_map<std::string, SEXP>` string-interning cache, which requires no Python registration for normal use. For the remaining three (`eval`, `findVar`, `findVarInFrame`), the guides specify a function-pointer bridge: Python registers a callback

at startup, and the C code calls it when needed.

PROTECT/UNPROTECT. Without a garbage collector, the GC protection protocol reduces to no-ops. SEXP lifetime is managed by the Python caller via explicit cleanup after the `.Call` boundary returns.

The header architecture. The guides collectively specify a layered DAG of fake headers, from `fake_arena.h` (lowest level, arena allocator) through type definitions (`INTSXP.h`, `SEXP.h`), accessors (`INTEGER.h`, `REAL.h`), memory management (`R_alloc.h`, `R_Free.h`), I/O (`Rprintf.h`, `error.h`), and finally the Category E function-pointer bridges (`R_getVar.h`, `install.h`, `eval.h`).

5 Phase 4: Fake C++ Header Generation and Compilation Bug Audit

5.1 Motivation

With implementation blueprints established in Phase 3, this phase translates each Markdown guide into actual compilable C++ code. Generating correct headers for 51 interdependent items is non-trivial: ODR violations, include-guard mismatches, and linkage conflicts are easy to introduce when headers are generated item by item without a global view. The phase therefore includes a systematic post-generation bug audit.

5.2 Header Generation

The `/generate-r-extern-fake-headers` (see Appendix D.1) skill was invoked:

```
/generate-r-extern-fake-headers rpart/src/ r_extern_analysis/combined.csv \  
r_extern_analysis/fake_guides/ r2py_rpart/r_fake_headers/
```

The `@generate-r-extern-fake-header` sub-agent was invoked sequentially for each of the 51 items, producing the corresponding `.h` file in `r2py_rpart/r_fake_headers/`. The generation order preserved the dependency-safe first-occurrence order from the CSV. After all 51 items, the master entry-point header `fake_R.h` was assembled, including all 51 headers in the correct dependency order beginning with `fake_arena.h` as the mandatory first include.

The arena allocator `fake_arena.h` was authored as a prerequisite: it provides `ArenaFrame` (a RAII guard for `.Call` boundaries), `arena_alloc(bytes)` (allocates from the current active frame), and `arena_calloc(n, size)`. Every `.Call` boundary wrapper that calls `R_alloc/ALLOC` must declare `ArenaFrame _frame`; as its first local variable; the destructor walks the frame's linked list of allocated blocks and frees them, then restores the previous frame pointer. The allocator is `thread_local` to support the Python ctypes model where each `.Call` invocation runs on a single thread.

5.3 Post-Generation Bug Audit

Following generation, a systematic compilation bug audit was performed. The primary tool was a pair of grep commands over all 53 headers:

```
grep -rh "#define_FAKE_" *.h | sort -u
grep -rh "#ifndef_FAKE_" *.h | sort -u
```

Seven compilation bugs were identified and fixed. The root causes fell into three classes: (1) include guard name mismatches (a header that first defines a symbol does not set the guard that a later header checks before emitting its fallback); (2) linkage conflicts (**static inline** fallbacks in companion headers conflicting with **inline** definitions in primary headers); and (3) multi-TU correctness for function-pointer variables.

The most consequential bug was that the four Category E function pointer variables (`g_install_fn`, `g_findVar_fn`, `g_findVarInFrame_fn`, `g_eval_fn`) were declared **static** at namespace scope in their header files. In C++, a **static** variable in a header has internal linkage: every translation unit that includes the header gets its own independent copy. In a multi-translation-unit shared library build, Python’s `register_install_fn(fn)` call would update only the copy belonging to the translation unit containing the registration function; all other units would retain `nullptr` and throw an error when the callback path is exercised. The fix was to change all four from **static** to **inline**, which gives C++17 shared-variable semantics: the linker selects one definition across all translation units.

The complete list of bugs fixed is summarized in Table 1.

Table 1: Compilation bugs identified and fixed during the Phase 4 audit.

Bug	Files Modified	Root Cause and Fix
1	<code>error.h</code>	<code>Rf_warning</code> defined without setting guard; <code>warning.h</code> fallback fired; both definitions in same TU. Fixed by adding <code>#define FAKE_RF_WARNING_DEFINED</code> in <code>error.h</code> .
2	<code>R_getVar.h</code> , <code>eval.h</code>	Function pointer variables declared static ; each TU got its own copy; Python registration only updated one. Fixed by changing to inline .

Continued on next page

Bug	Files Modified	Root Cause and Fix
3	DL_FUNC.h, two stubs	DL_FUNC.h used guard FAKE_R_BOOLEAN_H; Rboolean.h used FAKE_RBOOLEAN_DEFINED; mismatch caused Rboolean to be defined twice. Fixed by unifying guard names.
4	PROTECT.h	PROTECT.h used two separate guards per function; SEXP.h used one unified guard; mismatch caused R_ProtectWithIndex/R_Reprotect to be redefined (PROTECT_INDEX is typedef int, so types are identical—pure guard error). Fixed by unifying guards in PROTECT.h.
5	error.h	error.h defined Rf_error without any guard; R_NilValue.h (earlier in include chain) had already defined it with guard set. Fixed by wrapping error.h's definition in the matching guard.
6	DL_FUNC.h	Registration functions defined as inline in DL_FUNC.h without setting their individual guards; companion headers emitted static inline fallbacks; inline/static conflict is ill-formed in C++. Fixed by adding three #define FAKE_R*_DEFINED lines in DL_FUNC.h.
7	DL_FUNC.h	DL_FUNC.h declared struct DllInfo {} without setting FAKE_DLLINFO_DEFINED; DllInfo.h then attempted typedef struct DllInfo_fake DllInfo, redefining DllInfo as a different type. Fixed by adding the guard definition.

After all fixes, the 53 headers in `r2py_rpart/r_fake_headers/` are correct with respect to the ODR, guard consistency, linkage consistency, and multi-TU correctness.

6 Phase 5: Pure-C Entry-Point Generation and Fake-Header Migration

6.1 Motivation

With compilable fake headers available, the remaining steps needed before Python can call the C code are: (1) generate wrapper functions that translate between the SEXP-based internal interface and a plain-array interface callable without R, and (2) replace the real R API `#include` directives in the C source files with `#include "fake_R.h"`.

The entry-point wrappers are necessary because the internal C functions take SEXP arguments (opaque R list structures) and return SEXP values. Python cannot construct SEXP objects directly. The wrappers accept plain `int*/double*` arrays (matching how `.Call()` in R pre-converts its arguments) and perform the SEXP construction and deconstruction internally, returning results as plain arrays that Python can read.

6.2 Phase 5A: Entry-Point Generation

The `/generate-r-extern-raw-entry-points` (see Appendix E.1) skill was invoked:

```
/generate-r-extern-raw-entry-points rpart/R/ r2py_rpart/src/ \  
  r2py_rpart/r_fake_headers/ r2py_rpart/c_entry_points/
```

The skill first scanned all `.R` files in `rpart/R/` for `.Call()` and `.External()` invocations. Five invocations were found across four files; no `.External()` invocations exist. For each of the five C functions (`pred_rpart`, `init_rpcallback`, `rpartexp2`, `xpred`, `rpart`), the `@generate-r-extern-raw-entry-point` sub-agent generated a corresponding `*_c.c` file in `r2py_rpart/c_entry_points/`. All five wrappers follow a common structural convention:

- `#include "fake_R.h"` is placed before the `extern "C" {` block to avoid C++ template-in-linkage-spec errors.
- `ArenaFrame _frame;` is declared as the first local variable to own all arena allocations made during the call.
- An exception boundary is implemented via `noexcept` plus catch blocks for `RError`, `std::bad_alloc`, `std::exception`, and the catch-all (...), each of which writes a null-terminated error string to the caller-supplied error buffer.

The `init_rpcallback_c.c` wrapper is the most architecturally complex: the `rho`, `expr1`, and `expr2` arguments (R environments and expressions) must be passed as opaque `void*` handles, since they represent R interpreter state that cannot be expressed as plain arrays. This is the point where the Python-side frame registry (described in Phase 9) must construct fake SEXP objects wrapping numpy buffers and register the Category E function-pointer callbacks before the C function is called.

6.3 Phase 5B: Bulk Fake-Header Migration

The `/modify-c-folder-with-fake-headers` (see Appendix F.1) skill was invoked:

```
/modify-c-folder-with-fake-headers r2py_rpart/src/ r2py_rpart/c_entry_points/ \  
r2py_rpart/r_fake_headers/fake_R.h
```

All 35 `@modify-c-file-with-fake-headers` sub-agents were dispatched simultaneously (the modifications are fully independent). Of the 35 files, only 5 required modification: `rpart.h`, `init.c`, `pred_rpart.c`, `rpart_callback.c`, and `rpartexp2.c`. The remaining 30 files include only package-internal headers and receive their R API access transitively through `rpart.h`, which was itself modified. No R API symbols required suppression (comment-out) in any file—the fake header inventory from Phase 4 covered 100% of the R API surface.

Anonymous struct fix. During compilation verification, 14 files failed with:

```
rpart.h:76:3: error: '<unnamed>_struct>_rp', declared using unnamed type,  
is used but never defined [-fpermissive]
```

`rpart.h` declared the global `rp` struct as an anonymous type, valid in C but rejected by g++ in strict C++ mode. The fix changed `EXTERN struct {` to `EXTERN struct rp_globals_t {`, giving the type a tag name that satisfies C++ linkage requirements. This is the only code change introduced by Phase 5B that was not a mechanical header-include replacement.

After the fix, all 36 translation units (31 package sources + 5 entry-point wrappers) compiled without errors using:

```
g++ -std=c++17 -x c++ -I<fake_headers_dir> -I<c_folder> \  
-Wall -Wno-unused-variable -Wno-unused-parameter -fsyntax-only <file>
```

The 35 files that emitted warnings produced only two pre-existing benign warnings: a multi-line comment in `R_CallMethodDef.h`'s documentation block, and a semantically identical redefinition of `#define _(String) (String)` between `fake_R.h` and `rpart.h`.

7 Phase 6: R Package Structural Dependency Analysis

7.1 Motivation

Having established the C compilation infrastructure in Phases 1–5, the project turns to the R source layer. The 36 R source files define 47 functions; these functions must be converted to Python in a dependency-aware order (callee before caller) and each must be provided with the correct context about its internal and external dependencies. This phase produces the machine-readable structural analysis—JSON dependency maps and a topological level CSV—that drives the systematic conversion in Phase 9.

7.2 Methodology

The `/analyze-r-folder-dependencies` (see Appendix G.1) skill was invoked:

```
/analyze-r-folder-dependencies rpart/R/
```

The `@analyze-r-file-dependencies` sub-agent was applied sequentially to each of the 36 R source files. Each agent enumerated all function calls (including calls inside anonymous closures and nested functions) and classified each into one of three categories: *language dependencies* (base R or standard library functions), *internal dependencies* (functions defined in sibling `.R` files), or *external dependencies* (`.Call()`/`.External()` invocations targeting compiled C routines). Results were saved to `structural_analysis/R/<stem>.json`.

The pre-existing `dependency_graph.py` and `dependency_levels.py` scripts were then run to produce the graph visualization and the topological level CSV `structural_analysis/dependency_levels.`

7.3 Key Findings

The full call graph contains 262 nodes (36 file nodes, 49 function nodes, 172 language-dependency nodes, 5 external-dependency nodes) and 745 edges. Only 5 of the 36 files contain `.Call()` invocations, confirming that the R-to-C interface is narrow.

The 47 functions are stratified into 6 dependency levels. The two deepest utilities are `compress` (level 5, the subtree-overlap elimination algorithm inside `rpartco`) and `tree.depth` (level 5, computes node depth from binary node number via `floor(log2(node))`). The most widely referenced internal functions are `formatg` (called by all four method initialisers and `labels.rpart`), `na.rpart` (called by `model.frame.rpart`, `rpart`, and `xpred.rpart`), and `rpart.control`, `rpart.matrix`, and `rpartcallback` (each called by both `rpart` and `xpred.rpart`).

The `rpart` and `xpred.rpart` functions are the two primary orchestrators, each pulling in `na.rpart`, `rpart.matrix`, `rpartcallback`, and all four method-dispatch targets dynamically via `get(paste("rpart", method, sep="."))`.

8 Phase 7: Python Build Infrastructure

8.1 Motivation

The C compilation infrastructure from Phases 4–5 must be integrated into a proper Python packaging workflow before any Python-level testing or distribution is possible. This phase establishes a working `pip install` flow for the `r2py_rpart` package and designs the CI/CD pipeline for multi-platform wheel publication on PyPI. The reference design for the CI/CD infrastructure is the analogous project `r2py_kernsmooth` at `python-KernSmooth/r2py_kernsmooth/`.

8.2 Build System Bugs and Fixes

Bug 1: C++17 math built-in undefined references. `meson.build` originally compiled all sources with `-std=c++17`. On the project's GCC 11/glibc system, C++17 mode causes `<cmath>` to declare `std::iscanonical`, `std::issignaling`, `std::iseqsig`, and related overloads that reference GCC built-ins absent from this system's glibc. A manual test confirmed that `-std=c++14` produces zero undefined references. The fix was to change `default_options: ['cpp_std=c++17']` to `['cpp_std=c++14']` and update the `c_args` accordingly. The `inline` variable declarations in the fake headers (a C++17 feature) continue to compile under GCC 11 in C++14 mode as an accepted language extension.

Bug 2: Math library not linked. After the C++14 fix, `log`, `sqrt`, `pow`, and related standard math functions produced undefined references at link time. In C++14 mode these resolve to real library calls requiring `-lm`. The fix added `link_args: ['-lstdc++', '-lm']` to `meson.build`.

Bug 3: `_find_lib()` searched the source tree. The `_find_lib()` function in `__init__.py` computed its search path from `__file__`, which points into the source directory when `pytest` runs from the source tree rather than a site-packages install. The fix extended the function with a two-stage search: first the package directory co-located with `__file__`, then the `platlib/purelib` site-packages directories obtained via `sysconfig.get_path()`.

Bug 4: `_check_error()` received a numpy array. All public wrapper functions passed the error buffer as a numpy uint8 array to `_check_error()`, which called `ffi.string()` on it. `cffi`'s `ffi.string()` requires a `_CDataBase` object. The fix added a branch that handles `numpy.ndarray` input by splitting the byte sequence at the first null terminator.

8.3 CI/CD Pipeline

A GitHub Actions workflow `.github/workflows/python-publish.yml` was designed to build wheels for five platforms: Linux x86_64, Linux aarch64, macOS Intel (x86_64), macOS Apple Silicon (arm64), and Windows x86_64. The workflow uses `cibuildwheel` and is triggered by a GitHub Release publication event. It mirrors the `KernSmooth` structure with two adaptations. First, `rpart` requires no external system library (no BLAS, no Fortran runtime), so the Linux and macOS `before-all` hooks that install `openblas` are omitted. Second, on macOS, Apple clang links against `libc++` rather than GCC's `libstdc++`; `meson.build` was updated with a conditional:

```
if host_machine.system() == 'darwin'
    cxx_stdlib_link = ['-lc++']
else
    cxx_stdlib_link = ['-lstdc++']
endif
```

After all four bug fixes, the package builds successfully with `pip install -no-build-isolation .` and the existing `rpartexp2` tests pass.

9 Phase 8: Language Dependency Analysis and Conversion Guide Generation

9.1 Phase 8.1: Language Dependency Extraction

9.1.1 Motivation

Converting an R function to Python requires not just knowledge of the function’s own logic but also precise documentation of how each base-R language construct it uses should be translated. Base-R functions like `tapply`, `unlist`, `match`, `lapply`, and `model.frame` have non-trivial semantics that do not map one-to-one to any single Python or NumPy operation. Before conversion guides can be generated for these constructs, their usage locations across the entire codebase must be inventoried to identify all structurally distinct call patterns.

9.1.2 Methodology

The `/extract-r-folder-language-dependencies` (see Appendix H.1) skill was invoked:

```
/extract-r-folder-language-dependencies rpart/R/ structural_analysis/R/
```

Each of the 36 R files was processed by the `@extract-r-file-language-dependencies` sub-agent, which was provided both the R source file and its corresponding structural analysis JSON from `structural_analysis/R/`. The agent identified all invocation sites of the language dependencies catalogued in the JSON, extracted the line number and call body for each, and wrote a 4-column CSV to `language_dependency_analysis/R/`. The 36 CSVs were merged by the script `language_dependency_analysis/combine_csvs.py` into `combined_table.csv`.

9.1.3 CSV Encoding Defect

The initial merge failed with a `pandas.errors.ParserError`. A diagnostic loop identified `rpart.anova.csv` as the offending file: `call_body` fields containing R string literals (which use `"` as delimiters) were written with bare `"` characters rather than the RFC 4180-compliant doubled `""` escape sequence, causing the CSV parser to split fields at interior double-quotes. The file was corrected by rewriting it with Python’s `csv.writer` using `quoting=csv.QUOTE_MINIMAL`.

9.1.4 Key Findings

The combined table contains 1,298 invocation records across 172 distinct language dependencies. The most frequently called base-R functions are `stop` (80 call sites), `c` (67), `length` (63), `is.null` (60), `cat` (41), and `match` (35). The `rpartcallback.R` file is the densest single-function file: 84 rows reflecting heavy use of `stop` (14 calls), `length` (13 calls), and `assign` (9 calls) for environment construction.

9.2 Phase 8.2: Conversion Guide Generation

9.2.1 Motivation

With the full inventory of 172 language dependencies and their usage contexts available, conversion guides can be generated. Each guide documents the R semantics of one language construct, identifies all structurally distinct call patterns found in the rpart codebase, and provides concrete Python (NumPy/pandas/stdlib) equivalents with notes on type coercion, indexing conventions, and edge cases. These guides serve as the primary reference for the `@convert-r-function-to-python` agents in Phase 9.

9.2.2 Methodology

The `/generate-language-dependency-conversion-guides` (see Appendix I.1) skill was invoked:

```
/generate-language-dependency-conversion-guides rpart/R/  
language_dependency_analysis/combined_table.csv
```

The 172 unique `language_dependency` values were extracted from the combined table using Python's `csv.DictReader` (required because `awk` mishandles multi-line quoted fields). Per-dependency CSV subsets were written to a session scratchpad. Dependency names containing filesystem-unsafe characters were encoded with a custom scheme (e.g., `names<->names_LT_-.md`). The 172 `@generate-language-dependency-conversion-guide` agents were dispatched across 9 parallel batches of up to 20 agents each, all as background tasks. Output guides were written to `language_dependency_analysis/conversion_guides/`.

9.2.3 Verification Edge Case

Two guides are dot-prefixed (`.getXlevels.md` and `.checkMFCclasses.md`) and therefore invisible to shell globbing (`ls *.md` or `glob('*.md')`). Final verification required `os.listdir` or `find` to obtain the correct count of 172. An artifact directory `.tmp_csv/` was also present and correctly excluded from the count.

9.2.4 Representative Guide Complexity

The guides cover a wide range of translation complexity. Simple constructs like `any` map directly to `np.any()`. Complex constructs require pattern-specific guidance: `tapply` requires three distinct translations depending on whether the grouping is a free-form factor, a factor with explicit levels for zero-filling, or dense integer bin indices. `t` (transpose) requires distinguishing between a pure orientation transpose (`.T` numpy view) and a transpose-for-C-call that must be followed by `.ravel()` to produce a column-major flat array. `unlist` requires four distinct patterns depending on whether the input is a list of arrays, a list with NULL slots, a column-major DataFrame, or a named scalar/array packing for a `.Call` argument.

10 Phase 9: R-to-Python Function Conversion and Package Assembly

10.1 Phase 9.1: Batch R-to-Python Function Conversion

10.1.1 Motivation

With all prerequisite artifacts in place—structural dependency maps (Phase 6), language dependency conversion guides (Phase 8.2), and a working C back-end (Phases 4–5)—the systematic translation of the 36 R source files and 47 functions to Python can begin. The conversion must respect the topological order established by the dependency graph: functions whose output is consumed by other functions must be converted first.

10.1.2 Methodology

The `/convert-r-folder-to-python` (see Appendix J.1) skill was invoked:

```
/convert-r-folder-to-python rpart/R/ structural_analysis/R/  
    structural_analysis/dependency_levels.csv  
    language_dependency_analysis/conversion_guides/ conversion_results/R/
```

The skill parsed the dependency level CSV and sorted the 36 files in descending order of their maximum dependency level, ensuring that deep callees were converted before their callers. Within each multi-function file, functions were further sorted leaf-to-root. For example, `rpartco.R` was processed with `compress` (level 5) before `rpartco` (level 4); `rpart.exp.R` with `drate2` (level 2) before `rpart.exp` (level 1). Independent leaf functions across different files were dispatched in parallel.

Each `@convert-r-function-to-python` sub-agent was provided the R source file, the structural dependency JSON, the full set of 172 conversion guides, and the dependency level CSV. The output for each function was a JSON file with keys `imports`, `function_prototype`, and `function_body`, saved to `conversion_results/R/<stem>/<function>.json`.

10.1.3 Cross-Cutting Translation Decisions

The following R-to-Python idiom mappings were applied uniformly across all conversions:

R Construct	Python Equivalent
R 1-based indexing	Python 0-based indexing throughout
<code>list(a=1, b=2)</code> (named list)	<code>dict</code>
<code>missing(arg)</code>	Module-level <code>_MISSING = object()</code> sentinel
<code>match(x, table, 0L)</code>	Dict-based lookup; 0 or -1 for not found
<code>rep(x, n)</code>	<code>[x] * n</code> or <code>np.repeat</code>
<code>rbind/cbind</code>	<code>np.vstack/np.column_stack</code> or <code>np.hstack</code>
<code>cumsum(c(...))</code>	<code>np.cumsum(np.concatenate([...]))</code>
<code>pmax/pmin</code>	<code>np.maximum/np.minimum</code>
<code>tapply(wt, factor(y), sum)</code>	<code>pd.Series.groupby(pd.Categorical).sum()</code>
Column-major matrix for <code>.Call</code>	<code>np.reshape(..., order='F')</code>
<code>t(Y)</code> before <code>.Call</code>	<code>Y.T.ravel()</code>
<code>.Call("C_rpart")</code>	<code>from r2py_rpart import _rpart_c</code>
<code>new.env()/assign()/get()</code>	Python <code>dict</code>
<code>asNamespace("rpart") + environment<-</code>	No-op; omitted
<code>UseMethod(...)</code>	Dispatch stub with <code>NotImplementedError</code>
R graphics (<code>plot</code> , <code>polygon</code> , ...)	matplotlib equivalents

10.1.4 Notable Per-File Decisions

Several functions required non-mechanical decisions. `compress` in `rpartco.R` is a self-recursive nested function; it was extracted as a standalone module-level function with the closure variables (`is_leaf`, `nspc`) passed as explicit parameters. Similarly, `drate2` (nested inside `rpart.exp`) and `tfun` (nested inside `rpart`) were extracted as standalone helpers and also redefined inline in their enclosing function's conversion.

The largest function, `rpart()` (300 lines), required careful handling of the ordered-factor categorical split matrix construction, classification probability vector normalization, and assembly of the final output dictionary. R's S3 class assignment `class(ans) <- "rpart"` was replaced with a sentinel entry `ans["_rpart_class"] = "rpart"`, and `attr()`-stored metadata was handled via `ans["_xlevels"]` and `ans["_ylevels"]`.

The `eval.parent(model.frame(...))` pattern—which dynamically constructs a model frame from the parent call object—was identified as a hard boundary that cannot be faithfully reproduced in Python without a formula parser. All such sites were stubbed with `NotImplementedError`, to be filled in once patsy integration is available. The interactive `snip.rpart.mouse` function was also stubbed, as terminal interaction is not reproducible in Python.

10.2 Phase 9.2: Package Assembly, API Restructuring, and Test Validation

10.2.1 Motivation

The 47 JSON function definitions produced in Phase 9.1 exist as isolated per-function artifacts. This phase assembles them into importable Python modules, resolves inter-module imports, establishes the package’s public API, and validates the assembled package against a test suite.

10.2.2 Package Assembly

The `/combine-python-functions-into-folder` (see Appendix K.1) skill was invoked:

```
/combine-python-functions-into-folder conversion_results/R/  
r2py_rpart/r2py_rpart/ r2py_rpart/r2py_rpart/
```

A Python script read each conversion subdirectory’s JSON files, deduplicated imports, prepended the `from __future__ import annotations` and `from typing import Any` boilerplate, and assembled functions as prototype-plus-body blocks. All 36 files were written successfully. Twenty-six of the 36 output filenames contained dots from R naming conventions (e.g., `labels.rpart.py`) and were renamed by replacing internal dots with underscores (e.g., `labels_rpart.py`) to make them importable as Python modules.

Four categories of import-time errors were identified and fixed before tests could run:

1. **Undefined `_MISSING` sentinel** in four files: `rpart.py`, `rpart_control.py`, `summary_rpart.py`, `xpred_rpart.py`. The `_MISSING` sentinel is used as a default argument value (evaluated at function definition time); it was defined in some files but absent in others. Fixed by adding `_MISSING = object()` before each affected function.
2. **Stale absolute imports** in `rpart_exp.py` and `rpart_poisson.py`: the conversion agent emitted paths of the form `from conversion_results.R.formatg.R.formatg import formatg`, which are unreachable at runtime. Fixed by replacing with relative imports: `from .formatg import formatg`.
3. **Module-level import `patsy`** in `rpart_matrix.py`: `patsy` is not installed in the target environment. The import was moved inside the function body as a lazy import.
4. **Missing sibling-module imports in `rpart.py`**: the converted `rpart()` function called `rpart_anova`, `rpart_class`, `rpart_control`, `rpart_exp`, `rpart_matrix`, `rpart_poisson`, `rpartcallback`, and `importance` without importing them (R’s namespace made them available automatically). Eight relative imports were added.

10.2.3 API Restructuring

The original `__init__.py` exposed the four C-level wrapper functions as `rpart`, `pred_rpart`, `xpred`, and `rpartexp2`. After assembly, the high-level Python interfaces from the converted modules claimed the same names. The C wrappers were renamed to private symbols (`_rpart_c`, `_pred_rpart_c`, `_xpred_c`, `_rpartexp2_c`), and the high-level Python functions

were promoted to the public API via `__all__`. The three modules that previously imported the C wrappers by their old public names were updated accordingly.

10.2.4 Runtime Bugs in `rpart.py`

Three bugs in the assembled `rpart.py` were discovered and fixed during test development:

1. **Off-by-one errors in `isplit[:, 0]` variable indices.** The C entry point `rpart_c` returns `isplit[:, 0]` as 1-based variable indices (R convention), where index 0 is the sentinel for leaf nodes and index $k \geq 1$ refers to the k -th predictor. The conversion had incorrectly treated these indices as 0-based in three locations: the split row-name construction, the `isord` array lookup, and the `frame["var"]` column. All three were corrected by subtracting 1 before 0-based array accesses.
2. **parms dict serialization for classification method.** `rpart_class` returns `init["parms"]` as a dict with keys `prior`, `loss`, and `split`. The original serialization attempted `np.array(init["parms"]).ravel()`, which produces a 0-dimensional object array on a dict and raises `TypeError`. The fix iterates over `.values()` and concatenates the flattened arrays.
3. **Circular import avoidance.** `rpart.py` imports `_rpart_c` lazily inside the function body (not at module level) to avoid a circular import: `__init__.py` imports `rpart` from `rpart.py` at module level; a module-level import of `_rpart_c` from `r2py_rpart` would trigger `__init__.py` to run again before it has finished loading.

10.2.5 Build System Fix

After reinstalling with `pip install -no-build-isolation .`, the test runner reported `ModuleNotFoundError: No module named 'r2py_rpart.formatg'`. Inspection showed that only `__init__.py` was present in site-packages. The root cause is that `meson-python` requires every `.py` file to be enumerated explicitly in `py.install_sources()`; it does not auto-discover Python source files. The original listing had only one entry. The fix added all 36 new modules, bringing the total to 37 entries. This is a structural property of `meson-python` that differs from `setuptools` (`find_packages()`) and must be maintained: any session that adds new Python modules must also update `meson.build` in the same change.

10.2.6 Test Results

A test suite of 11 tests for the primary `rpart()` function was written and all 11 pass. The tests cover: output dictionary structure, `frame` type, `cptable` column names, `where` vector semantics, method recording for both `anova` and `class` methods, splitting behavior on a synthetic signal-vs-constant dataset, and error handling for invalid inputs.

10.3 Phase 9.3: R Interpreter Item Bridge and cffi Migration

10.3.1 Motivation

Two gaps remained after Phase 9.2 before the package could support the user-defined split path (`method=4`). First, the C side was missing five helper functions needed by the Python-side frame registry to construct fake SEXP objects wrapping numpy buffers (`make_real_sexp`, `make_int_sexp`, `make_env_sexp`, `free_sexp_helper`, and access to `get_make_sexp_error`). Second, the converted `rpartcallback.py` used the `ctypes` module for callback registration and SEXP construction, which is incompatible with the package’s `cffi`-based `_lib` object.

10.3.2 Methodology

The `/add-r-interpreter-items` (see Appendix L.1) skill was invoked:

```
/add-r-interpreter-items r_extern_analysis/fake_guides/  
r2py_rpart/c_entry_points/ r2py_rpart/r2py_rpart/ r2py_rpart/
```

The skill identified three Category E guides whose opening blockquote begins with **R Interpreter Item**. (strict prefix match): `eval.md`, `findVar.md`, and `findVarInFrame.md`. The `install.md` guide was excluded because its opening blockquote notes that the `install` function is best-effort fakeable in C++ without a Python function pointer—its built-in hash-map implementation in `install.h` is sufficient for normal use.

The `@add-r-interpreter-items` agent was invoked once with the full content of all relevant files, including the fake guides, the existing `init_rpcallback.c.c` wrapper (authoritative protocol reference), `fake_R.h`, `__init__.py`, and `rpartcallback.py`.

10.3.3 C Helper File

A new file `r2py_rpart/c_entry_points/interpreter_helpers.c.c` was written exporting five `extern "C" noexcept` functions:

Function	Signature	Purpose
<code>make_real_sexp</code>	<code>void*(void*, int)</code>	Allocate REALSXP node; data not copied
<code>make_int_sexp</code>	<code>void*(void*, int)</code>	Allocate INTSXP node; data not copied
<code>make_env_sexp</code>	<code>void*(void)</code>	Allocate minimal ENVSXP identity shell
<code>free_sexp_helper</code>	<code>void(void*)</code>	Free SEXPREC node only (not data buffer)
<code>get_make_sexp_error</code>	<code>const char*(void)</code>	Return thread-local error string

The function `call_install` was *not* added as a new C-level function: `install.h` already emits it as a weak (W) `extern "C" inline` symbol via `-fkeep-inline-functions`, and redefining it in a `.c` file would cause a duplicate-symbol linker error. The correct approach is to declare it in `ffi.cdef()` only, reusing the pre-existing weak symbol.

10.3.4 rpartcallback.py Rewrite

The full rewrite replaced every `ctypes` FFI call with its `cffi` equivalent:

Old (ctypes)	New (cffi)
<code>ctypes.CFUNCTYPE(...)</code>	<code>ffi.callback(..., fn)</code>
<code>class _SEXPREC(ctypes.Structure): ...</code>	<code>_lib.make_real_sexp(buf, size)</code>
<code>(ctypes.c_char * 1)()</code>	<code>ffi.new("char[1]")</code>
<code>ctypes.cast(node, ctypes.c_void_p).value</code>	<code>int(ffi.cast("uintptr_t", node))</code>
<code>arr.ctypes.data_as(ctypes.c_void_p)</code>	<code>ffi.from_buffer(arr)</code>
<code>_lib.register_*_fn.restype = ...</code>	<code>_lib.register_*_fn(ffi.cast("void *", cb))</code>

A subtle lifetime consideration: `cffi`'s `ffi.from_buffer(arr)` returns a `cdata` object that pins the numpy buffer for the duration of the call, but both the buffer object and the `from_buffer` result must be kept in the `_cffi_keep_alive` list to prevent premature garbage collection. The SEXP node pointers allocated by `make_real_sexp` for the `yback/wback/xback/nback` buffers are kept conservatively in the same list, even though after `init_rpcallback_c` returns the C static globals point directly into the numpy buffers, not into the SEXP nodes.

A distinction that was explicitly preserved: the three helper functions `_iptr`, `_dptr`, `_cptr` in `__init__.py` call `arr.ctypes.data` (numpy's `.ctypes` attribute, which returns an integer memory address) and pass it to `ffi.cast(...)`. This is valid `cffi` usage—`numpy.ndarray.ctypes` is a numpy feature unrelated to the Python `ctypes` module—and was correctly retained without modification.

10.3.5 Results

After the build (`pip install -no-build-isolation .`), symbol verification confirmed the presence of all expected symbols in `_rpart_core.so`: `make_real_sexp`, `make_int_sexp`, `make_env_sexp`, `free_sexp_helper` as strong (T) symbols from `interpreter_helpers_c.c`, and `call_install` as a weak (W) symbol from `install.h`. All 11 regression tests continue to pass. The `method=4` user-defined splits code path is architecturally complete at the FFI layer; end-to-end testing with actual user-defined split functions has not yet been performed.

A Phase 1: C Dependency Analysis

A.1 Skill: /analyze-c-folder-dependencies

```
---
name: analyze-c-folder-dependencies
description: Analyzes all C files in a folder to extract structural dependency
             information file by file, categorizing calls strictly into internal and external
             dependencies.
---
```



```
# Analyze a C Folder for Dependencies

## Description

When provided with a target folder, your task is to recursively scan for all C files (
    including the header files), use the '@analyze-c-file-dependencies' agent on each
    file individually, and save the results as structured JSON files.

## Execution Steps

### Step 1: Scan for C Files

Recursively scan the provided target folder and all of its subdirectories for files
    with '.c' or '.h' extensions. Compile a complete list of their absolute file paths.

### Step 2: Process and Analyze

Iterate through the list of identified C files sequentially. For each file:
1. Invoke the '@analyze-c-file-dependencies' agent against the file and record its
    final JSON output.
2. If the command fails or throws an error for a specific file, log the error to the
    console and immediately continue to the next file in your list. Do not halt the
    entire operation.

### Step 3: Save Output

For every successfully analyzed C file, save the resulting JSON output according to
    these strict rules:
* **Naming Convention:** Use the exact name of the original C file, but append the
    extension '.json' (e.g., 'init.c' becomes 'init.c.json').
* **Output Location:** If the user did not specify the output folder, save the new '.
    json' file in 'c_refactor_analysis/{folder_name}/', where the '{folder_name}' is
    the name of the target folder. The relative path must be the same as its
    corresponding original C file.
```

A.2 Agent: @analyze-c-file-dependencies

```
---
name: analyze-c-file-dependencies
description: Analyzes a provided C file to extract structural dependency information
    function by function (including function-like macros), categorizing calls strictly
    into internal and external dependencies.
---

# Analyze a C File for Dependencies

## Description

When provided with a C file ('.c' or '.h'), your task is to parse the file, identify
    every user-defined function or function-like macro, and trace all function-like
    calls (functions and function-like macros) made within their bodies. You must
    categorize these dependencies into strict groups and output the result as a
    structured JSON object.
```

Execution Steps

Step 1: Extract Function or Macro Declarations

Scan the target C file and identify every user-defined function or function-like macro defined within it (e.g., 'int my_func() { ... }'). Create a list of these function or macro identifiers.

Step 2: Analyze Function or Macro Bodies and Trace Calls

For every function or macro identified in Step 1, thoroughly analyze its execution body. Extract the identifier of every function or function-like macro invoked within that body.

Exclusions:

- * Completely ignore functions or macros that are part of the C standard library (e.g., 'printf', 'malloc', 'strlen', 'size_t').
- * Extract **only the identifier** (e.g., 'PyArg_ParseTuple', '.Call'), never the arguments or the parentheses.

Categorize the extracted identifiers into the following two mutually exclusive lists. You may need to traverse and scan other files in the same project folder according to the included header files, and you have to use the following heuristics:

* **Internal Dependencies:**

* **Definition:** Identifiers for functions or macros defined within the same C file, or those that appear in other '.c' or '.h' files within the same project folder, which are included via '#include' directives into the target file.

* **External Dependencies:**

* **Definition:** Identifiers for functions or macros belonging to known external, third-party libraries or language bindings (for example, R API calls like 'Rf_protect', Python C API calls like 'Py_Initialize', etc.). Only if you cannot find the definition of the function or macro in the same file or any other file in the same project folder, you should classify it as an external dependency.

Step 3: Generate JSON Output

Construct a JSON object where each key is the name of a function or macro identified in Step 1. The value for each key must be an object containing exactly two arrays of strings: 'internal_dependencies' and 'external_dependencies'. Deduplicate all identifiers within these arrays.

Output Format Schema

You must output valid JSON strictly adhering to this structure. Do not include markdown code blocks if the system expects raw JSON.

```
{
  "my_custom_algorithm": {
    "internal_dependencies": ["helper_function_within_file", "project_utils_init"],
    "external_dependencies": ["Rf_protect", "PyArg_ParseTuple"]
  },
  "helper_function_within_file": {
```

```

    "internal_dependencies": [],
    "external_dependencies": []
  }
}

```

B Phase 2: R External Item Extraction

B.1 Skill: /extract-c-folder-r-extern-items

```

---
name: extract-c-folder-r-extern-items
description: Recursively processes a folder of C files to generate detailed CSV reports
  of R external references (macros, types, functions).
---

# Extract R External Items from a C Folder

## Description

When provided with a target directory containing C source or header files, your task is
to process every C file (‘.c’ or ‘.h’) recursively. You will invoke the ‘@extract-
c-file-r-extern-items’ agent on each identified C file and save the extracted
dependency data strictly as ‘.csv’ files in a mirrored output directory.

## Execution Steps

### Step 1: Scan for C Source and Header Files

Recursively scan the provided target directory and all of its subdirectories. Compile a
complete list of absolute file paths for every file with a ‘.c’ or ‘.h’ extension.

### Step 2: Sequential Processing and Error Isolation

Iterate through the compiled list of identified C files sequentially. For each file:
1. Execute File-Level Agent: Invoke the ‘@extract-c-file-r-extern-items’ agent,
  providing it with the current ‘.c’ or ‘.h’ file path as its target.
2. Error Handling: If the sub-agent fails, times out, or throws an error for a
  specific file, log a clear warning to the console containing the file’s path.
  Immediately proceed to the next C file in the list. Do not halt the entire batch
  execution due to individual file failures.

### Step 3: Save Output as CSV

For every successfully analyzed C file, save the tabular string output returned by the
sub-agent according to these strict rules:
* Data Format: The output content must be saved strictly as a flat CSV file
  matching the schema defined by the sub-agent.
* Naming Convention: Use the exact base name of the original C file, but append it
  with ‘.csv’ (e.g., ‘rpart.c’ becomes ‘rpart.c.csv’ and ‘rpart.h’ becomes ‘rpart.h.
  csv’).
* Output Directory Mapping: Maintain the original directory hierarchy. Save the ‘.
  csv’ file into the user-specified output directory, mirroring the exact relative

```

```
path of the source C file relative to the input target folder.
* **Default Output Directory:** If the user did not explicitly specify an output
  directory, create a default directory path named 'r_extern_analysis/{
  target_folder_name}/' (where '{target_folder_name}' is the base name of the target
  folder being scanned) and mirror the relative paths there.
```

B.2 Agent: @extract-c-file-r-extern-items

```
---
name: extract-c-file-r-extern-items
description: Extracts detailed location and usage data for R external references (
  macros, types, functions) in a C file.
---

# Extract R External Items from a C File

## Description

When provided with a C source file ('.c' or '.h'), your task is to locate every
instance where an R external item is referenced. An R external item is defined as
any type identifier, struct, macro, or function name that is declared in a header
file located in the folder '~/.conda/envs/r-to-python/lib/R/include/'.

You must extract the exact line numbers, the containing code statements, and the
specific R header file where each item is declared. The final output must be a
strictly formatted, sorted CSV.

## Execution Steps

### Step 1: Identify Non-Local C Identifiers

Parse the provided C file to identify all macros, type definitions, structs, and
function calls. Exclude any identifiers that are clearly defined within the
provided C file itself or are part of the standard C library (e.g., 'stdio.h', '
stdlib.h' elements).

### Step 2: Cross-Reference with R Headers

For the identifiers flagged in Step 1, use your search tools (e.g., 'grep' or file-
reading capabilities) to search both the C files that the current C file includes
and the directory '~/.conda/envs/r-to-python/lib/R/include/'. Identify which of
these items are declared in those R header files in '~/.conda/envs/r-to-python/lib/
R/include/'.

For every matched R external item, record the following six data points:
* **'external_item'**: The exact name of the referenced item (e.g., 'SEXP', 'PROTECT',
  'ALLOC').
* **'header_file'**: The specific header file within '~/.conda/envs/r-to-python/lib/R/
  include/' where the item is declared (e.g., 'R.h', 'Rinternals.h').
* **'category'**: A string flag indicating the category of the item, which can only be
  among "type", "variable", "function". Types include those defined by typedefs,
  structs, enums, or type-like macros. Variables are identifiers that represent
  values. Functions are identifiers that are invoked with parentheses '()', which can
```

```

    be either functions or function-like macros. Note that there is no category for
    macros since they are always crafted to mimic the behavior of the three categories
    above.
* **'file_name'**: The name of the provided '.c' or '.h' file you are currently
    analyzing.
* **'line_number'**: The exact line number in the analyzed file where the statement
    containing the reference begins.
* **'context_statement'**: The complete code snippet containing the reference.
    * If the statement spans multiple lines, capture the entire complete statement from
    start to the terminating semicolon ';', or the complete function prototype/macro
    definition.
    * Consolidate multi-line statements into a single string by escaping internal line
    breaks or replacing them with spaces.

### Step 3: Format and Sort Output

Compile the extracted data points into a CSV format. Sort the output primarily by '
    line_number' (ascending), and secondarily by 'external_item' (alphabetically).

## Output Format Schema

You must output valid CSV data strictly adhering to the structure below. Include the
    exact header row. **CRITICAL:** You must enclose the 'context_statement' strings in
    double quotes (") to prevent commas or formatting elements within the C code
    from breaking the CSV structure.

'''csv
external_item,header_file,category,file_name,line_number,context_statement
SEXP,R.h,type,rpart.c,40,"SEXP x = PROTECT(ALLOC(REALSXP, n));"
PROTECT,R.h,function,rpart.c,40,"SEXP x = PROTECT(ALLOC(REALSXP, n));"
ALLOC,R.h,function,rpart.c,40,"SEXP x = PROTECT(ALLOC(REALSXP, n));"
REALSXP,Rinternals.h,type,rpart.c,40,"SEXP x = PROTECT(ALLOC(REALSXP, n));"
'''

```

C Phase 3: Fake Header Implementation Guides

C.1 Skill: /generate-r-extern-fake-guides

```

---
name: generate-r-extern-fake-guides
description: Orchestrates the batch generation of fake C++ header implementation guides
    for R's C API external items, enabling R package C source files to compile and run
    without linking against libR.so.
---

# Generate R External Item Fake Header Guides

## Description

When provided with a target 'base_folder', a CSV file matching the schema below, and a
    target 'output_directory', your task is to orchestrate the batch processing of this
    dataset. You must segment the CSV by unique 'external_item' values and iteratively

```

invoke the '@generate-r-extern-fake-guide' agent for each subset to generate the resulting guides to the specified output directory.

The goal of each guide is to document how to implement a drop-in C++ fake for a given R C API external item -- a type, macro, constant, or function -- so that the original package C source files can be compiled and linked without 'libR.so', and called directly from Python.

Expected CSV Schema

Example:

```
'''csv
external_item,header_file,category,file_name,line_number,context_statement
DL_FUNC,R_ext/Rdynload.h,type,init.c,12,"static const R_CallMethodDef CallEntries[] = {
    {"init_rpcallback", (DL_FUNC) &init_rpcallback, 5}, {"rpart", (DL_FUNC) &rpart
, 11}, {"xpred", (DL_FUNC) &xpred, 15}, {"rpartexp2", (DL_FUNC) &rpartexp2, 2},
    {"pred_rpart", (DL_FUNC) &pred_rpart, 12}, {NULL, NULL, 0} };"
DllInfo,R_ext/Rdynload.h,type,init.c,23,R_init_rpart(DllInfo * dll)
SEXP,Rinternals.h,type,rpart.c,41,"SEXP rpart(SEXP ncat2, SEXP method2, SEXP opt2, SEXP
    parms2, SEXP xvals2, SEXP xgrp2, SEXP ymat2, SEXP xmat2, SEXP wt2, SEXP ny2, SEXP
    cost2)"
INTSXP,Rinternals.h,type,pred_rpart.c,139,"SEXP where = PROTECT(allocVector(INTSXP, n))
;"
REALSXP,Rinternals.h,type,rpart.c,241,"cptable3 = PROTECT(allocMatrix(REALSXP, xvals >
    1 ? 5 : 3, rp.num_unique_cp));"
PROTECT,Rinternals.h,macro,rpart.c,194,"which3 = PROTECT(allocVector(INTSXP, n));"
allocVector,Rinternals.h,function,rpart.c,194,"which3 = PROTECT(allocVector(INTSXP, n))
;"
error,R_ext/Error.h,function,rpart.c,203,"error("%s", errmsg);"
R_alloc,R_ext/Memory.h,function,rpart.c,123,"rp.xdata = (double **) ALLOC(rp.nvar,
    sizeof(double *));"
'''
```

Execution Steps

Step 1: Identify Unique Dependencies

Read and parse the input CSV file. Extract a definitive list of all unique values present in the 'external_item' column (e.g., 'DL_FUNC', 'SEXP', 'INTSXP', 'PROTECT', 'allocVector', 'error', 'R_alloc').

Note: Assume the CSV is pre-sorted. Identical external items are grouped sequentially to allow for efficient linear processing.

Step 2: Extract Segments and Invoke Sub-agents

Iterate sequentially through the list of unique 'external_item' values one after another. Sequential execution is mandatory here: foundational fake definitions (e.g., the 'SEXP' struct, the 'SEXPTYPE' enum, the arena allocator) must be generated before the items that depend on them (e.g., 'allocVector', 'INTEGER', 'REAL'), and each newly generated guide may be referenced by the agent processing the next item. For each unique item, strictly execute the following sub-steps:

1. ****Extract the CSV Subset:**** Isolate all rows where the ‘external_item’ column matches the current target. Prepend the exact CSV header row (‘external_item, header_file,category,file_name,line_number,context_statement’) to this subset to create a valid, standalone CSV-formatted string.
2. ****Execute Fake Guide Agent:**** Invoke the ‘@generate-r-extern-fake-guide’ agent. Pass the target ‘base_folder’, the newly created CSV subset string, and the specified output directory as inputs.
3. ****Strict Error Handling:**** If the agent execution fails, times out, or throws an error at any point, you must:
 - Log a precise error message to the console: ‘ERROR: Failed to generate fake guide for {external_item}. Proceeding to next item.’
 - Immediately continue the loop with the next unique ‘external_item’ in the list. Do not halt or abort the overall batch execution.

C.2 Agent: @generate-r-extern-fake-guide

```

---
name: generate-r-extern-fake-guide
description: Generates a fake C++ header implementation guide for a specific R C API
  external item, enabling the original R package C source files to compile and run
  without linking against libR.so.
---

# Generate an R External Item Fake Header Guide

## Description

When provided with a target ‘base_folder’ and a CSV text input, your task is to
generate a comprehensive fake header implementation guide for the specified ‘
external_item’ (e.g., ‘SEXP’, ‘INTSXP’, ‘allocVector’, ‘error’, ‘R_alloc’). This
guide documents how to implement a drop-in C++ replacement for that item so that
the original package C source files compile and link without ‘libR.so’, and can be
called directly from Python.

The input CSV will strictly adhere to the following example schema:
“‘csv
external_item,header_file,category,file_name,line_number,context_statement
allocVector,Rinternals.h,function,rpart.c,194,"which3 = PROTECT(allocVector(INTSXP, n))
;"
allocVector,Rinternals.h,function,rpart.c,241,"cptable3 = PROTECT(allocMatrix(REALSXP,
xvals > 1 ? 5 : 3, rp.num_unique_cp));"
allocVector,Rinternals.h,function,rpart.c,261,"dnode3 = PROTECT(allocMatrix(REALSXP,
nodecount, (3 + rp.num_resp));"
allocVector,Rinternals.h,function,rpartexp2.c,47,"SEXP keep = PROTECT(allocVector(
INTSXP, n));"
“‘

The generated guide must:
1. Define the ‘external_item’'s core functionality within R's C API.
2. Analyze its specific usage within the provided CSV dataset, incorporating
  surrounding source code context.
3. Provide a complete C++ fake implementation strategy governed by the three invariants
  below, including concrete, syntax-highlighted C++ code for every distinct usage

```

pattern found in the CSV.

Three Invariants That All Fake Implementations Must Respect

These invariants apply to every guide generated by this agent, regardless of the specific 'external_item'. They are non-negotiable design constraints:

****Invariant 1 -- C++ error and warning style.****

All R error and warning mechanisms ('Rf_error', 'Rf_warning', and their '#define' aliases 'error', 'warning') must be replaced by pure C++ equivalents. Errors ('Rf_error') must throw a C++ exception -- either 'std::runtime_error' or a project-defined subclass 'RError : public std::runtime_error'. Warnings ('Rf_warning') must write to 'stderr' via 'std::fprintf(stderr, "Warning: ...")'. Under no circumstances may 'longjmp', 'setjmp', R condition handlers, or 'abort()' be used.

****Invariant 2 -- Arena-based memory management without modifying original code.****

All R memory allocation functions used in the package C source ('R_alloc', 'R_chk_calloc', 'R_chk_realloc', 'R_chk_free', 'S_alloc', and the macros 'ALLOC', 'CALLOC', 'R_Calloc', 'R_Realloc', 'R_Free', 'Memcpy') must be reimplemented using a thread-local arena allocator. The arena is structured as a stack of frames: each frame corresponds to one invocation of a top-level '.Call'-style C function. An 'ArenaFrame' RAII guard is pushed at function entry and popped (freeing all frame allocations) at function exit. The original C source files are ****not modified**** -- the arena is completely transparent to them because the same function names and macro names are preserved in the fake headers.

The canonical arena interface that all memory-related guides must reference is:

```
““cpp
// fake_arena.hpp (foundational; generated once and referenced by all memory guides)
#include <cstdlib>
#include <cstring>
#include <stdexcept>
#include <vector>

struct ArenaBlock {
    char    *data;
    size_t  used;
    size_t  capacity;
};

struct Arena {
    std::vector<ArenaBlock> blocks;
    static constexpr size_t kDefaultBlockSize = 65536;

    void *alloc(size_t bytes) {
        bytes = (bytes + 7) & ~size_t(7); // 8-byte align
        if (blocks.empty() || blocks.back().used + bytes > blocks.back().capacity) {
            size_t cap = bytes > kDefaultBlockSize ? bytes : kDefaultBlockSize;
            ArenaBlock b;
            b.data      = static_cast<char *>(std::malloc(cap));
            if (!b.data) throw std::bad_alloc();
            b.used      = 0;
            b.capacity  = cap;
        }
    }
};
```



```

        blocks.push_back(b);
    }
    char *ptr = blocks.back().data + blocks.back().used;
    blocks.back().used += bytes;
    return ptr;
}

void reset() {
    for (auto &b : blocks) std::free(b.data);
    blocks.clear();
}

~Arena() { reset(); }
};

// One arena per thread; stack discipline maps to nested .Call invocations
inline thread_local std::vector<Arena> gArenaStack;

struct ArenaFrame {
    ArenaFrame() { gArenaStack.emplace_back(); }
    ~ArenaFrame() { gArenaStack.back().reset(); gArenaStack.pop_back(); }
};

inline void *arena_alloc(size_t bytes) {
    return gArenaStack.back().alloc(bytes);
}

inline void *arena_calloc(size_t n, size_t size) {
    void *p = arena_alloc(n * size);
    std::memset(p, 0, n * size);
    return p;
}
'''

**Invariant 3 -- Best-effort faking of R interpreter items.**
Some R C API items ('eval', 'findVar', 'findVarInFrame', 'install', 'R_getVar', '
    R_GlobalEnv', 'R_NaString', environments 'ENVSXP', and expression objects 'LANGSXP
    '/'EXPRSXP') require a running R interpreter and cannot be fully faked. For these
    items, the guide must:
- Clearly mark the item as an **R Interpreter Item** at the top of the guide.
- Explain exactly why a complete fake is impossible.
- Provide the best-effort Python interop alternative: a C function pointer that Python
    registers via 'ctypes.CFUNCTYPE' before calling the enclosing '.Call' function. The
    fake implementation calls this function pointer instead of 'eval(expr, rho)' or '
    findVar(sym, rho)'. The guide must show the full C++ stub (the function pointer
    declaration, the registration function, and the stub body) and the corresponding
    Python-side 'ctypes' registration code.
- Note which specific package methods or code paths require the item (e.g., "method=4
    user-defined splits in rpart") and state that those paths remain unavailable unless
    the Python function pointer is registered.

## Execution Steps

### Step 1: Ingest Inputs and Gather Deterministic Context

```

- ****Path Resolution:**** Iterate through each row in the CSV. Construct the absolute file path by concatenating the provided 'base_folder' with the 'file_name'.
- ****File Inspection:**** Navigate to the exact 'line_number' in each target file. Read at least 15 lines above and 15 lines below the target line. Use the chosen window to understand: the C types of all arguments and return values, the scope and lifetime of variables, how the target 'external_item' interacts with adjacent R API calls (especially 'PROTECT'/'UNPROTECT', 'ALLOC', and SEXP accessors), and what data flows into and out of the item.
- ****Header Resolution:**** You must read the local header file at '~/conda/envs/r-to-python/lib/R/include/{header_file}' whenever the 'external_item's struct definition, macro expansion, or function signature is not explicitly clear from the source code window. Also check '~/conda/envs/r-to-python/lib/R/include/R_ext/' for sub-headers included transitively. If a header is not found in either 'base_folder' or '~/conda/envs/r-to-python/lib/R/include/', explicitly state that the header is missing and proceed with the information available. Never search other locations.

Step 2: Resolve External Dependencies

If local headers do not provide complete definitions, consult the following references in order:

1. '.Call/.External' documentation: '<https://search.r-project.org/R/refmans/base/html/CallExternal.html>'
2. Writing R Extensions (memory and error chapters): '<https://cran.r-project.org/doc/manuals/R-exts.html>'
3. Package-specific CRAN documentation (e.g., '<https://cran.r-project.org/web/packages/rpart/>').

***Requirement:** Before writing the fake implementation, search the output directory for any previously generated fake guides. If a guide already exists for a type or function that the current 'external_item' depends on (e.g., 'allocVector' depends on 'SEXP' and 'SEXPTYPE'; 'INTEGER' depends on 'SEXP'; 'error' depends on 'RError'), read that guide and ensure the new fake implementation is consistent with -- and references -- the already-established fake definitions.

Step 3: Design the Fake C++ Implementation

Classify the 'external_item' into exactly one of the five implementation categories below. Apply the rules for that category to design the fake.

****Category A -- Type or Enum Constant**** (e.g., 'SEXP', 'SEXPTYPE', 'INTSXP', 'DL_FUNC', 'DllInfo', 'Rboolean', 'R_len_t', 'Rbyte'):

- Provide a complete C++ struct, typedef, or 'enum' definition that matches the observable behavior of the R type in the package source.
- For 'SEXP' specifically: define 'SEXP' as a C++ struct with at minimum a 'SEXPTYPE' type, 'int length', 'int nrow', 'int ncol', and a 'void *data' field. 'SEXP' is 'SEXP *'. SEXP objects are heap-allocated and not garbage-collected in the fake runtime; the caller owns their lifetime.
- For 'DllInfo' and 'R_CallMethodDef': define as minimal structs or no-op placeholders sufficient for 'init.c' to compile; the registration functions are no-ops.

****Category B -- Accessor Macro or Inline Function**** (e.g., 'INTEGER', 'REAL', 'LOGICAL', 'RAW', 'COMPLEX', 'asInteger', 'asReal', 'asLogical', 'LENGTH', 'XLENGTH', 'nrows', 'ncols', 'TYPEOF', 'STRING_ELT', 'VECTOR_ELT', 'CHAR', 'PRINTNAME', 'R_CHAR

```

    ):
- Implement as a C++ 'inline' function (preferred) or as a '#define' macro if the
  original is a macro with no safe inline equivalent.
- The implementation casts 'sexp->data' to the appropriate pointer type and returns it
  (for vector accessors) or extracts a scalar (for 'asInteger', 'asReal').
- For 'nrows' and 'ncols': return 'sexp->nrow' and 'sexp->ncol'.

**Category C -- Allocation or Memory Function** (e.g., 'allocVector', 'allocMatrix', '
  R_alloc', 'R_chk_calloc', 'R_chk_realloc', 'R_chk_free', 'R_Calloc', 'R_Free', '
  Malloc', 'mkChar', 'PROTECT', 'UNPROTECT'):
- All heap memory obtained by SEXP-allocating functions ('allocVector', 'allocMatrix',
  'mkChar') must use 'std::malloc'/'new' and must set 'sexp->type', 'sexp->length', '
  sexp->nrow', 'sexp->ncol', and 'sexp->data' correctly.
- All arena-allocated functions ('R_alloc', 'R_chk_calloc', 'R_chk_realloc', 'S_alloc')
  must delegate to 'arena_alloc' / 'arena_calloc' from 'fake_arena.hpp' (Invariant
  2).
- 'R_chk_free' and 'R_Free' free heap memory (for non-arena allocations). They must NOT
  free arena memory; the arena does that at frame destruction.
- 'PROTECT' and 'UNPROTECT' are **no-ops** in the fake runtime (there is no garbage
  collector). Implement as empty inline functions or identity macros. '
  PROTECT_WITH_INDEX' and 'REPROTECT' are also no-ops.
- For every allocation function, explicitly show how 'ArenaFrame' must be declared at
  the entry of each top-level '.Call' function to ensure cleanup (referencing
  Invariant 2).

**Category D -- Error, Warning, or Print Function** (e.g., 'Rf_error', 'Rf_warning', '
  Rprintf', 'REprintf'):
- Apply Invariant 1 strictly.
- 'Rf_error(fmt, ...)': format the message with 'vsprintf' into a 'std::string', then
  'throw RError(msg)' where 'RError' is defined as 'struct RError : public std::
  runtime_error { using std::runtime_error::runtime_error; };'. The '#define error
  Rf_error' alias from 'R_ext/Error.h' still works unchanged.
- 'Rf_warning(fmt, ...)': format and write to 'stderr' via 'std::fprintf'. Do not throw
  .
- 'Rprintf(fmt, ...)': delegate to 'std::printf'.
- 'REprintf(fmt, ...)': delegate to 'std::fprintf(stderr, ...)' .
- Show how the '.Call' boundary wrapper (in Python's ctypes glue) must catch 'RError'
  and translate it into a Python exception via a 'try { ... } catch (const RError &e)
  { /* set error flag */ }' block.

**Category E -- R Interpreter Item** (e.g., 'eval', 'findVar', 'findVarInFrame', '
  R_getVar', 'install', 'R_GlobalEnv', 'R_NaString', 'ENVXP', 'LANGXP'):
- Apply Invariant 3 strictly.
- Mark the item as an R Interpreter Item in the guide header.
- Provide the full C++ function pointer stub: global function pointer variable, a C-
  linkage registration function callable from Python, and the stub body that calls
  the pointer (or throws 'RError' if the pointer has not been registered).
- Provide the complete Python-side 'ctypes' snippet showing type declaration, callback
  implementation, and registration call.
- State which code paths in the package require this item and which paths do not (so
  the user knows exactly what is blocked and what is free).

## Output Format Schema

```

You must output the final fake header guide strictly in Markdown format. Use the exact headings and structure outlined below:

1. Overview of ‘{external_item}’ in R API

*Provide a 2-3 sentence definition of the external item: what it is, its role in R’s C API, its typical inputs and outputs, and -- if it falls under Invariant 3 -- a clear statement that it is an **R Interpreter Item**.*

2. Contextual Usage Analysis

Summarize the data extracted from Step 1. List: the C types of all arguments and return values seen in context, which other R API items appear alongside this one (e.g., ‘PROTECT’, ‘ALLOC’, adjacent accessor calls), and the distinct implementation patterns present across the CSV rows. Identify which patterns share the same fake strategy and which require separate treatment.

3. Fake C++ Implementation Strategy

State which implementation Category (A-E) this item belongs to. Then describe:

- *The specific fake mechanism chosen and why it satisfies the invariants.*
- *For Category C items: which allocations go to the arena vs. the heap, and how the ‘ArenaFrame’ guard interacts with the function that calls this item.*
- *For Category D items: the ‘RError’ exception class, how ‘vsprintf’ is used to format variadic arguments, and what the ‘.Call’ boundary ‘try/catch’ looks like.*
- *For Category E items: why a full fake is impossible, what the function pointer bridge achieves, and the exact boundary at which Python registers the callback.*
- *Any ‘#define’ aliases from the original header (e.g., ‘#define error Rf_error’, ‘#define PROTECT(s) Rf_protect(s)’) that must be preserved in the fake header so that the original source files compile unchanged.*

4. Fake Implementation Examples

For each distinct usage pattern identified in Step 2, create a subsection formatted exactly as follows:

Pattern: [Brief Description of Pattern, e.g., Allocate 1D Integer Vector]

- **Locations:** [List ‘file_name:line_number’ for every CSV row belonging to this pattern]
- **Original R API Usage:** [A representative code snippet from the source, showing the original SEXP-based call]
- **C++ Fake Implementation:** [A complete, syntax-highlighted C++ snippet providing the fake definition for this item, in the context of this pattern. Include the ‘ArenaFrame’ guard if the pattern involves arena memory, and the ‘try/catch’ boundary wrapper if the pattern involves error throwing.]

```

- **Arena / Memory Notes:** [Only if this pattern allocates memory. Specify: does this
  allocation go to the arena or to the heap? When is it freed? What happens if the
  allocation fails?]
- **Python Interop Notes:** [Only for Category E items. Show the full ctypes
  registration snippet and the C++ function pointer stub. Explain which Python-side
  data structures map to which C arguments.]
- **Explanation:** [Detail the mechanical changes: what the fake header defines, what
  ‘#define’ aliases are preserved, and why the original source file compiles without
  modification.]

---

### 5. Integration Requirements

*List every other ‘external_item’ whose fake guide must be included before this one in
the final fake header (i.e., compile-order dependencies). For each dependency,
state which specific definition from that guide is needed here (e.g., “‘SEXP’ guide
required -- provides the ‘SEXP_PREC’ struct used in ‘sexp->data’”). If there are no
dependencies, state “None.”*

---

## Output Saving Instructions

- **File Naming:** Save the generated guide as ‘{external_item}.md’ (e.g., ‘allocVector
  .md’, ‘INTSEXP.md’, ‘error.md’).
- **Output Directory:** Save the file to the user-specified output directory. If the
  user did not explicitly provide one, create and use the default directory path ‘r-
  extern-analysis/fake_guides/’.

```

D Phase 4: Fake Header Generation

D.1 Skill: /generate-r-extern-fake-headers

```

---
name: generate-r-extern-fake-headers
description: Generates fake C++ header files for all R C API external items listed in a
  CSV table, using pre-generated Markdown fake guides as implementation blueprints,
  and assembles a master entry-point header that replaces R.h and Rinternals.h for
  standalone builds.
---

# Generate Fake C++ Header Files for R C API External Items

## Description

Given four inputs -- a C source ‘base_folder’, a ‘csv_file’ mapping all R external
  items that need faking, a ‘guides_folder’ containing the pre-generated Markdown
  fake implementation guides (one per ‘external_item’), and an ‘output_folder’ for
  the resulting ‘.h’ files -- your task is to:

1. Parse the CSV to determine the ordered list of unique ‘external_item’ values.

```

2. Sequentially invoke the '@generate-r-extern-fake-header' agent for each item to produce a corresponding '{external_item}.h' file in the 'output_folder'.
3. After all individual header files have been generated, write a master entry-point header ('fake_R.h') in the 'output_folder' that '#include's every generated header in the correct dependency order, so that package C source files need only add one include directive to compile without 'libR.so'.

Expected CSV Schema

The 'csv_file' must adhere to the following example schema:

```
'''csv
external_item,header_file,category,file_name,line_number,context_statement
DL_FUNC,R_ext/Rdynload.h,type,init.c,12,"static const R_CallMethodDef CallEntries[] = {
    {"init_rpcallback", (DL_FUNC) &init_rpcallback, 5}, {"rpart", (DL_FUNC) &rpart
    , 11}, {"xpred", (DL_FUNC) &xpred, 15}, {"rpartexp2", (DL_FUNC) &rpartexp2, 2},
    {"pred_rpart", (DL_FUNC) &pred_rpart, 12}, {NULL, NULL, 0} };"
DllInfo,R_ext/Rdynload.h,type,init.c,23,R_init_rpart(DllInfo * dll)
INTSXP,Rinternals.h,type,pred_rpart.c,139,"SEXP where = PROTECT(allocVector(INTSXP, n))
;"
INTSXP,Rinternals.h,type,rpart.c,194,"which3 = PROTECT(allocVector(INTSXP, n));"
INTSXP,Rinternals.h,type,rpart.c,278,"inode3 = PROTECT(allocMatrix(INTSXP, nodecount,
6));"
INTSXP,Rinternals.h,type,rpart.c,285,"isplit3 = PROTECT(allocMatrix(INTSXP, splitcount,
3));"
INTSXP,Rinternals.h,type,rpart.c,293,"csplit3 = PROTECT(allocMatrix(INTSXP, catcount,
maxcat));"
INTSXP,Rinternals.h,type,rpartexp2.c,47,"SEXP keep = PROTECT(allocVector(INTSXP, n));"
REALSXP,Rinternals.h,type,rpart.c,241,"cptable3 = PROTECT(allocMatrix(REALSXP, xvals >
1 ? 5 : 3, rp.num_unique_cp));"
REALSXP,Rinternals.h,type,rpart.c,261,"dnode3 = PROTECT(allocMatrix(REALSXP, nodecount,
(3 + rp.num_resp));"
REALSXP,Rinternals.h,type,rpart.c,269,"dsplit3 = PROTECT(allocMatrix(REALSXP,
splitcount, 3));"
REALSXP,Rinternals.h,type,xpred.c,209,"predict2 = PROTECT(allocVector(REALSXP, n * ncp
* nresp));"
'''
```

Execution Steps

Step 1: Identify and Order Unique External Items

Read and parse the 'csv_file'. Extract the definitive ordered list of all unique values in the 'external_item' column, preserving their first-occurrence order within the file.

Step 2: Generate Individual Fake Header Files

Iterate sequentially through the ordered list from Step 1. Sequential execution is mandatory: earlier items (e.g., 'SEXP', 'SEXPTYPE', 'fake_arena') define types and memory structures that later items (e.g., 'allocVector', 'INTEGER', 'REAL', 'error') must reference. Each item's generated header must exist and be readable before the next agent is invoked. For each unique 'external_item', strictly execute the following sub-steps:

1. ****Extract the CSV Subset:**** Isolate all rows where the 'external_item' column matches the current target. Prepend the exact CSV header row ('external_item, header_file, category, file_name, line_number, context_statement') to form a valid, standalone CSV-formatted string.
2. ****Invoke the Header Generation Agent:**** Call the '@generate-r-extern-fake-header' agent, passing the 'base_folder', the CSV subset string, the 'guides_folder', and the 'output_folder'.
3. ****Non-Blocking Error Handling:**** If the agent fails, times out, or produces a syntactically invalid header, you must:
 - Log the precise error to the console: 'ERROR: Failed to generate fake header for {external_item}. Proceeding to next item.'
 - Immediately continue to the next 'external_item'. Do not halt the overall batch.

Step 3: Assemble the Master Entry-Point Header

After all individual agents have completed (including any that were skipped due to errors), assemble a master header file named 'fake_R.h' in the 'output_folder'. This file serves as the single drop-in replacement for R's standard includes ('R.h', 'Rinternals.h', 'Rdefines.h', 'R_ext/Error.h', etc.) in the package source files.

To build 'fake_R.h':

1. ****Scan the output folder**** for all '.h' files that were successfully generated in Step 2. Collect their filenames in the same order they were produced (which preserves the dependency-safe generation order from Step 1).
2. ****Write the master header**** with the following structure:
 - A file-level comment block identifying this as the fake R API entry point, listing the real R headers it replaces, and noting that it was auto-generated.
 - '#ifndef'/'#define'/'#endif' header guards and a '#pragma once' guard.
 - An '#include "fake_arena.h"' directive as the first include, since the arena underpins all memory allocation fakes. Add a comment noting that 'fake_arena.h' must be present in the same directory and is not auto-generated by this tool.
 - One '#include' line per successfully generated '{external_item}.h' file, in generation order.
 - A trailing comment block summarizing which original R headers are now covered and which 'external_item' values (if any) were skipped due to agent errors.
3. ****Do not include**** any '{external_item}.h' that was not successfully written to disk; only reference files that actually exist in the output folder.

D.2 Agent: @generate-r-extern-fake-header

```
---
name: generate-r-extern-fake-header
description: Writes a fake C++ header file for a specific R C API external item by
  translating the implementation blueprint from its pre-generated Markdown guide into
  compilable C++ code, so the original package source files build without libR.so.
---

# Generate a Fake C++ Header File for an R C API External Item

## Description
```

When provided with a ‘base_folder’, a CSV text snippet, a ‘guides_folder’ containing pre-generated Markdown fake implementation guides, and an ‘output_folder’, your task is to produce a single self-contained C++ header file -- ‘{external_item}.h’ -- that provides a drop-in fake implementation of the target R C API item. The original package C source files must be able to ‘#include’ this header (transitively through ‘fake_R.h’) and compile without any modification and without linking against ‘libR.so’.

The input CSV will strictly adhere to the following example schema:

```

““csv
external_item,header_file,category,file_name,line_number,context_statement
INTSXP,Rinternals.h,type,pred_rpart.c,139,"SEXP where = PROTECT(allocVector(INTSXP, n))
;"
INTSXP,Rinternals.h,type,rpart.c,194,"which3 = PROTECT(allocVector(INTSXP, n));"
INTSXP,Rinternals.h,type,rpart.c,278,"inode3 = PROTECT(allocMatrix(INTSXP, nodecount,
6));"
INTSXP,Rinternals.h,type,rpart.c,285,"isplit3 = PROTECT(allocMatrix(INTSXP, splitcount,
3));"
INTSXP,Rinternals.h,type,rpart.c,293,"csplit3 = PROTECT(allocMatrix(INTSXP, catcount,
maxcat));"
INTSXP,Rinternals.h,type,rpartexp2.c,47,"SEXP keep = PROTECT(allocVector(INTSXP, n));"
““

```

The generated header must:

1. Implement every function, macro, type, or constant that the ‘external_item’ represents, exactly as specified in the corresponding Markdown guide in ‘guides_folder’.
2. Preserve every ‘#define’ alias that the original R header defined for the item (e.g., ‘#define error Rf_error’, ‘#define PROTECT(s) Rf_protect(s)’) so that the original source files compile without any textual changes.
3. Satisfy the three invariants stated in the Markdown guide: C++ exception-based errors, arena-backed memory management without source modification, and best-effort Python function-pointer bridges for R interpreter items.

Execution Steps

Step 1: Ingest Inputs and Gather Deterministic Context

- ****Guide Ingestion:**** Read the Markdown guide at ‘{guides_folder}/{external_item}.md’. Extract and fully understand all five sections: Overview, Contextual Usage Analysis, Fake C++ Implementation Strategy, Fake Implementation Examples, and Integration Requirements. The guide is the authoritative blueprint; the steps below gather additional context to ensure the generated header is correct.
- ****Source File Inspection:**** Iterate through each row in the CSV. Construct the absolute source file path from ‘base_folder’ and the ‘file_name’ column. Navigate to the ‘line_number’ and read at least 15 lines above and below. Use these windows to confirm: the exact C types of arguments and return values that the item is used with, the scope and lifetime of surrounding variables, whether the item appears inside a function body or at file scope, and any adjacent R API calls that interact with it. If the source context reveals usage patterns not covered by the guide, note them explicitly in the generated header as code comments.
- ****Original Header Resolution:**** Read the original R header at ‘~/conda/envs/r-to-python/lib/R/include/{header_file}’ to extract the exact function signature, macro definition, or type layout declared there. Also check sub-headers under ‘~/conda/

envs/r-to-python/lib/R/include/R_ext/' for any transitively included definitions. If the header cannot be found in either 'base_folder' or the conda include path '~/conda/envs/r-to-python/lib/R/include/', state this in a comment within the generated '.h' file and proceed with the definition from the Markdown guide. Never search other locations.

Step 2: Resolve Header Dependencies

- ****Read Integration Requirements:**** Locate the "Integration Requirements" section (Section 5) of the Markdown guide. This section lists every other 'external_item' whose fake header must be included before the current one. For each listed dependency, verify that the corresponding '{dependency}.h' file exists in the 'output_folder'.
- ****Inspect Existing Headers:**** For each dependency that exists, read its content. Confirm that the types, structs, and function signatures it exposes are consistent with what the current item's guide expects. If a dependency is missing from 'output_folder', include a '// WARNING: dependency {dependency}.h not found in output_folder' comment in the generated header at the point of its '#include' directive and continue; do not abort.
- ****Check for Conflicts:**** Scan all existing headers in 'output_folder' to ensure the current item does not redefine a symbol already declared there. If a conflict exists, emit the new definition conditionally using '#ifndef' guards rather than duplicating it.

Step 3: Write the Fake Header File

Translate the Markdown guide's "Fake C++ Implementation Strategy" (Section 3) and "Fake Implementation Examples" (Section 4) into a complete, compilable C++ header file. Apply the following rules without exception:

****Rule 1 -- Use both '#ifndef' '#define' '#endif' header guards and '#pragma once'.** You have to add traditional '#ifndef'/'#define'/'#endif' include guards to prevent multiple inclusions. Also, place '#pragma once' as the very first non-comment line.

****Rule 2 -- Dependency includes come first.****

After the header guards, emit one '#include "{dependency}.h"' line per dependency identified in Step 2, in the order listed by the guide's Integration Requirements. If the item is in Category C (memory allocation), also emit '#include "fake_arena.h"' before all other includes.

****Rule 3 -- Implement exactly what the guide specifies for the item's category.****

- ****Category A (type or enum constant):**** Emit the C++ 'struct', 'typedef', or 'enum' definition. For 'SEXP'/'SEXPREF', output the full struct with 'SEXP_TYPE type', 'int length', 'int nrow', 'int ncol', and a 'union' of typed data pointers.
- ****Category B (accessor macro or inline function):**** Emit 'inline' C++ functions, not macros, for type safety. The function body casts 'sexp->data' or reads scalar fields from the 'SEXPREF' struct. Preserve any '#define' alias from the original header beneath the 'inline' definition.
- ****Category C (allocation or memory function):**** Implement heap-allocating functions ('allocVector', 'allocMatrix', 'mkChar') as 'inline' functions using 'new' / 'std::malloc'. Implement arena-delegating functions ('R_alloc', 'R_chk_calloc', 'S_alloc') as 'inline' functions calling 'arena_alloc' / 'arena_calloc' from 'fake_arena.h'. Implement 'PROTECT' and 'UNPROTECT' as no-op 'inline' functions. Place an 'ArenaFrame' usage comment block showing callers exactly where to declare '

```

    ArenaFrame _frame;' at their function entry.
- **Category D (error, warning, or print function):** Define the 'RError' struct
  exactly once (guard with '#ifndef FAKE_R_RERROR_DEFINED'). Implement 'Rf_error' as
  a '[[noreturn]]' 'inline' function that formats its variadic arguments into a 'std
  ::string' via 'vsprintf' and throws 'RError'. Implement 'Rf_warning' as an 'inline
  ' function that writes to 'stderr' via 'std::fprintf'. Implement 'Rprintf' and '
  Rfprintf' as 'inline' forwarding functions. Emit all '#define' aliases from the
  original R error header ('#define error Rf_error', '#define warning Rf_warning').
- **Category E (R interpreter item):** Emit the function pointer declaration as a
  global 'extern "C"' variable, the registration function with C linkage, and the
  stub body that calls the pointer or throws 'RError(...)' if the pointer is null.
  Wrap the entire block in a '// *** R INTERPRETER ITEM -- requires Python callback
  registration ***' comment block.

**Rule 4 -- Preserve all original R header '#define' aliases.**
After the implementation code, emit a block labelled '// --- Compatibility aliases (
  preserve original R API names) ---' containing every '#define' the original R
  header declared for this item. These ensure the unchanged package source files see
  the same names they used before (e.g., '#define allocVector Rf_allocVector' if the
  original header used that remapping).

**Rule 5 -- Emit an 'ArenaFrame' usage note for any Category C header.**
After the alias block, emit a comment block titled '// --- ArenaFrame usage ---' that
  shows the exact one-liner a caller must add at the top of any '.Call'-style
  function body that calls arena-backed functions defined in this header.

**Rule 6 -- No runtime R linkage.**
The generated header must introduce zero dependencies on 'libR.so'. Do not '#include'
  any original R headers. The only permitted system includes are from the c/C++
  standard library ('<stdlib.h>', '<stddef.h>', '<string.h>', '<stdexcept>', '<stdio.
  h>', '<stdarg.h>', '<vector>', '<string>', '<unordered_map>', etc.) and from other
  fake headers already in the 'output_folder'.

## Output Format Schema

The generated '.h' file must follow this exact structure, in order:

'''
// =====
// {external_item}.h -- Fake R API: {external_item}
// Original R header: {header_file}
// Category: {A|B|C|D|E} -- {category description}
// Auto-generated by generate-r-extern-fake-header agent.
// =====
#pragma once

// --- Standard library includes ---
#include <...>

// --- Fake R header dependencies ---
// (from Integration Requirements in {external_item}.md)
#include "fake_arena.h"          // if Category C
#include "{dependency_1}.h"
#include "{dependency_2}.h"

```

```
// --- R Interpreter Item warning block (Category E only) ---
// *** R INTERPRETER ITEM -- {external_item} ***
// A complete fake is impossible without an R interpreter.
// Register a Python callback via {external_item}_register() before use.
// Affected code paths: [list from guide]

// --- Core fake implementation ---
// [struct / typedef / enum / inline function / function pointer stub]
// One block per usage pattern from Section 4 of the Markdown guide.
// Each block is preceded by a comment citing the guide pattern name
// and the source file locations it covers.

// --- Compatibility aliases (preserve original R API names) ---
// [#define aliases from the original R header]

// --- ArenaFrame usage (Category C only) ---
// Declare 'ArenaFrame _frame;' as the first statement of any
// .Call-entry function that invokes arena-backed functions from
// this header. Example:
//   SEXP my_entry(SEXP x) { ArenaFrame _frame; ... }
'''
```

If the 'external_item' belongs to Category E (R interpreter item), the file must additionally contain, after the core fake implementation block, a Python-side usage note formatted as a multi-line C++ comment block that reproduces the full 'ctypes' registration snippet from the Markdown guide's "Python Interop Notes" section.

```
## Output Saving Instructions
- **File Naming:** Save the generated header as '{external_item}.h' (e.g., 'allocVector.h', 'INTSXP.h', 'error.h').
- **Output Directory:** Save the file to the user-specified 'output_folder'. If the user did not explicitly provide one, create and use the default directory 'r2py_rpart/r_fake_headers/'.
```

E Phase 5A: Pure-C Entry-Point Generation

E.1 Skill: /generate-r-extern-raw-entry-points

```
---
name: generate-r-extern-raw-entry-points
description: Scans R source files for .Call/.External invocations, identifies every C
             function called directly from R, and generates a pure C entry-point wrapper for
             each that replaces SEXP-based parameters and return values with plain int/double
             arrays callable from Python.
---

# Generate Pure C Entry Points for R-Callable C Functions

## Description
```

Given four inputs -- an 'r_base_folder' containing R source files, a 'c_base_folder' containing the package C source files, a 'fake_headers_folder' containing the pre-generated fake R C API headers, and an 'output_folder' for the resulting '.c' files -- your task is to:

1. Recursively scan every '.R' file in 'r_base_folder' to locate all '.Call' and '.External' invocations and record the C function name, call type, R call expression, and source location for each one.
2. Identify all unique C functions that are invoked directly from R, consolidating multiple call sites for the same function.
3. Sequentially invoke the '@generate-r-extern-raw-entry-point' agent for each unique C function to produce a '{c_function}_c.c' file in 'output_folder' containing a plain-C wrapper that exposes the function to Python without SEXP.

Execution Steps

Step 1: Scan R Source Files and Build the Call Site Table

Recursively scan 'r_base_folder' for all files with a '.R' or '.r' extension. For every R file found, parse it line by line to locate every '.Call(...)' and '.External(...)' invocation. For each invocation found, extract and record the following four data points:

- **'c_function'**: The name of the C function being called. This is the first argument to '.Call'/'External'. It may appear as:
 - A registered native symbol object (e.g., 'C_rpart', 'C_pred_rpart') -- strip the 'C_' prefix to obtain the C function name.
 - A quoted string literal (e.g., "rpart", "pred_rpart") -- use the string content directly.
- **'call_type'**: Either 'Call' (for '.Call(...)') or 'External' (for '.External(...)').
- **'r_file'**: The path of the R file, relative to 'r_base_folder'.
- **'r_line'**: The line number where the '.Call'/'External' expression begins.
- **'r_call_expression'**: The complete '.Call'/'External' expression, including all argument names or expressions passed after the function identifier. If the call spans multiple lines, consolidate it into a single string with internal whitespace collapsed.

Compile all records into an internal call site table adhering to this CSV schema:

```
'''csv
c_function,call_type,r_file,r_line,r_call_expression
rpart,Call,rpart.R,145,".Call(C_rpart, ncat, method, opt, parms, xvals, xgrp, ymat,
      xmat, wt, ny, cost)"
xpred,Call,xpred.rpart.R,52,".Call(C_xpred, ncat, method, opt, parms, xvals, xgrp, ymat
      , xmat, wt, ny, cost, all, cp, toprisk, nresp)"
pred_rpart,Call,predict.rpart.R,89,".Call(C_pred_rpart, dimx, nnode, nsplit, dimc, nnum
      , nodes2, vnum, split2, csplit2, usesur, xdata2, xmiss2)"
rpartexp2,Call,rpart.R,256,".Call(C_rpartexp2, dtimes, eps)"
init_rpcallback,Call,rpart.R,198,".Call(C_init_rpcallback, rho, ny, nr, expr1, expr2)"
'''
```

Step 2: Identify Unique C Functions and Consolidate Call Sites

From the call site table built in Step 1, extract the definitive ordered list of unique ‘c_function’ values, preserving the first-occurrence order across the scanned R files. For each unique C function, collect all rows from the table that share that ‘c_function’ value into a grouped CSV subset -- multiple rows exist when the same C function is called from more than one R file or location.

Step 3: Generate Entry-Point Files Sequentially

Iterate through the ordered list of unique C functions from Step 2 one at a time. For each function, strictly execute the following sub-steps:

1. ****Prepare the CSV Subset:**** Take all rows from the call site table where ‘c_function’ matches the current target. Prepend the exact header row (‘c_function,call_type,r_file,r_line,r_call_expression’) to form a valid, standalone CSV-formatted string.
2. ****Invoke the Entry-Point Agent:**** Call the ‘@generate-r-extern-raw-entry-point’ agent, passing the ‘r_base_folder’, ‘c_base_folder’, ‘fake_headers_folder’, ‘output_folder’, and the CSV subset string.
3. ****Non-Blocking Error Handling:**** If the agent fails, times out, or produces syntactically invalid C code, you must:
 - Log the precise error to the console: ‘ERROR: Failed to generate raw entry point for {c_function}. Proceeding to next function.’
 - Immediately continue to the next C function in the list. Do not halt the overall batch.

E.2 Agent: @generate-r-extern-raw-entry-point

```

---
name: generate-r-extern-raw-entry-point
description: Writes a pure C entry-point wrapper for a single R-callable C function,
             replacing its SEXP-based parameters and return value with plain int/double array
             arguments so the function can be called directly from Python without libR.so.
---

# Generate a Pure C Entry-Point Wrapper for an R-Callable C Function

## Description

When provided with an ‘r_base_folder’, a ‘c_base_folder’, a ‘fake_headers_folder’
containing pre-generated fake R C API headers, an ‘output_folder’, and a CSV text
snippet listing every R call site that invokes a specific C function, your task is
to produce a single C source file -- ‘{c_function}_c.c’ -- that wraps the original
SEXP-based function in a new ‘{c_function}_c’ entry point that Python can call
directly via ‘ctypes’ or ‘cffi’ using only plain ‘int’, ‘double’, and ‘char’
pointer arguments.

The input CSV will strictly adhere to the following schema:
'''csv
c_function,call_type,r_file,r_line,r_call_expression
pred_rpart,Call,predict.rpart.R,89,".Call(C_pred_rpart, dimx, nnode, nsplit, dimc, nnum
    , nodes2, vnum, split2, csplit2, usesur, xdata2, xmiss2)"
pred_rpart,Call,predict.rpart.R,201,".Call(C_pred_rpart, dimx2, nnode, nsplit, dimc,
    nnum, nodes2, vnum, split2, csplit2, usesur, xdata2, xmiss2)"
'''

```

The generated file must:

1. Implement a '{c_function}_c' function whose signature uses only plain C types -- 'int *', 'double *', 'int' (for scalar dimensions), and 'char *' (for error output) -- with no SEXP anywhere in the signature.
2. Internally construct fake SEXP objects from the plain input arrays, call the original '{c_function}' using the fake R API from 'fake_headers_folder', unpack the returned SEXP into the caller-provided output arrays, and propagate any 'RError' exception as a null-terminated string via an 'error_out' parameter.
3. Compile without 'libR.so' by relying exclusively on the fake headers in 'fake_headers_folder'.

Execution Steps

Step 1: Ingest Inputs and Gather Deterministic Context

- ****R Call Site Inspection:**** For every row in the CSV, construct the absolute R file path from 'r_base_folder' and the 'r_file' column. Navigate to the 'r_line' and read at least 20 lines above and below. From this window, extract:
 - The declared R type of each argument passed to '.Call'/'External' ('https://search.r-project.org/R/refmans/base/html/CallExternal.html') (look for preceding 'as.integer', 'as.double', matrix constructions, 'dim()' calls, etc.).
 - Whether each argument is a scalar, a 1-D vector, or a 2-D matrix (established by how R prepares the argument before passing it to '.Call').
 - The names the R code uses for output components after the '.Call' returns (e.g., 'out\$which', 'out\$cptable'), to understand the return structure.
 - Whether 'call_type' is 'Call' or 'External'. For 'External', note that all arguments arrive as a single SEXP pairlist; document this difference prominently in the generated file header comment but still provide the best-effort flat-array entry point.
- ****C Function Definition Lookup:**** Search 'c_base_folder' recursively for the file that defines '{c_function}' (the function whose first token matches 'SEXP\s+{c_function}\s*\('). Read the complete function signature and the first 60 lines of the function body to determine:
 - The exact number and SEXP parameter names.
 - Which SEXP parameters are used as integer vectors ('INTEGER()'), real vectors ('REAL()'), integer matrices ('INTEGER()' + 'nrows()'/'ncols()'), real matrices, or scalar integers/reals ('asInteger()'/'asReal()').
 - The type and structure of the return SEXP (scalar, 1-D vector of a specific type, or a named list 'VECSXP'). For a list return, read far enough into the function body to identify every 'SET_VECTOR_ELT' and 'SET_STRING_ELT' call, recording the index, name string, and SEXP type of each component.
- ****Fake Header Inspection:**** Read '{fake_headers_folder}/fake_R.h' (or 'fake_R.hpp') to confirm the 'SEXPREC' struct layout and the names of the fake allocation, accessor, and error functions. Read the individual header files for the specific items used by '{c_function}' (e.g., 'INTSXP.h', 'REALSXP.h', 'allocVector.h', 'PROTECT.h', 'INTEGER.h', 'REAL.h', 'error.h', 'R_alloc.h'). This ensures the construction and extraction helpers in the generated file are consistent with the fake header definitions already established.

Step 2: Derive the Plain-C Signature

Using the context gathered in Step 1, produce a complete mapping from the original SEXP parameters and return value to their plain-C equivalents. Apply the following rules in order:

****Input parameter mapping:****

- SEXP used as 'asInteger(x)' (scalar int) -> one 'int' value parameter.
- SEXP used as 'asReal(x)' (scalar double) -> one 'double' value parameter.
- SEXP used as 'INTEGER(x)' only, with 'LENGTH(x)' but no 'nrows'/'ncols' -> 'int *{name}, int {name}_len'.
- SEXP used as 'REAL(x)' only, with 'LENGTH(x)' but no 'nrows'/'ncols' -> 'double *{name}, int {name}_len'.
- SEXP used as 'INTEGER(x)' with 'nrows(x)'/'ncols(x)' -> 'int *{name}, int {name}_nrow, int {name}_ncol'.
- SEXP used as 'REAL(x)' with 'nrows(x)'/'ncols(x)' -> 'double *{name}, int {name}_nrow, int {name}_ncol'.
- SEXP that is a Category E R interpreter item (environment, expression, language object) -> replace with a 'void *{name}_callback' function pointer of an appropriate 'typedef'-d type, with a comment explaining the Python registration requirement. See the corresponding guide in 'fake_headers_folder' for the pointer type definition.

****Return value mapping:****

- SEXP that is a 1-D integer vector -> one output parameter 'int *{name}_out' (pre-allocated by caller) and one 'int {name}_len' parameter carrying the expected length.
- SEXP that is a 1-D real vector -> 'double *{name}_out', 'int {name}_len'.
- SEXP that is a named list ('VECSXP') -> one pair '{type} *{component}_out' / dimension parameter(s) per list component, plus one 'int *has_{optional_component}' flag parameter for any component whose presence is conditional (e.g., 'csplit' in 'rpart'). Derive the component types and sizes from the 'SET_VECTOR_ELT' analysis in Step 1.
- In all cases, append two trailing parameters at the end of the signature: 'char *error_out' (a caller-allocated buffer that receives the error message if 'RError' is thrown) and 'int error_out_len' (its capacity in bytes). The presence of a non-empty 'error_out' after the call signals an error to the caller.

Document the complete derived signature in a structured comment block at the top of the generated file, listing each parameter, its direction ('in', 'out', 'in/out'), its plain-C type, and the original SEXP parameter or return component it corresponds to.

Step 3: Write the '{c_function}_c.c' File

Produce the complete C source file applying the following rules without exception:

****Rule 1 -- Includes.****

Begin with a file-level comment block (see Output Format Schema). Then emit:

```

'''_C
#include "fake_R.h"    /* or whichever master fake header exists in fake_headers_folder
    */
#include <stdlib.h>
#include <string.h>
#include <stdarg.h>
'''

```

Do not include any original R headers. Do not include the C source file that defines ‘{c_function}’; instead use an ‘extern’ declaration (Rule 2).

****Rule 2 -- Extern declaration of the original function.****

Emit the exact SEXP-based function prototype from the C source file, prefixed with ‘extern’, so the linker can resolve it from the compiled package object:

```
““C
extern SEXP {c_function}(SEXP param1, SEXP param2, ...);
““
```

****Rule 3 -- Static inline SEXP construction helpers.****

For each distinct plain-C -> SEXP conversion required by the input parameters (determined in Step 2), write a ‘static inline’ helper function that allocates a ‘SEXP’ on the heap, sets its ‘type’, ‘length’, ‘nrow’, ‘ncol’, and ‘data’ fields using the fake header definitions, and returns the ‘SEXP’. The helper must ****not**** copy the data -- it must point ‘data’ directly at the caller-supplied array, since the original C function only reads from these inputs. Name helpers descriptively: ‘make_int_vec’, ‘make_real_vec’, ‘make_int_mat’, ‘make_real_mat’, ‘make_int_scalar’, ‘make_real_scalar’. Each helper must be guarded against a null data pointer and throw ‘RError’ (from the fake error header) if ‘malloc’ fails.

****Rule 4 -- Static inline SEXP extraction helpers.****

For each distinct SEXP -> plain-C conversion required by the return value (determined in Step 2), write a ‘static inline’ helper that reads from a ‘SEXP’ using the accessor functions from the fake headers and copies into a caller-supplied output array. Name helpers descriptively: ‘extract_int_vec’, ‘extract_real_vec’, ‘extract_int_mat’, ‘extract_real_mat’. For list components, write one helper per component type, taking the list SEXP, the component index, and the output pointer. Each extraction helper must validate the SEXP type tag matches what is expected and call ‘Rf_error’ (which throws ‘RError’) if there is a mismatch.

****Rule 5 -- The ‘{c_function}_c’ entry-point function.****

Write the entry point with the exact plain-C signature derived in Step 2. Declare it ‘noexcept’ -- ‘ctypes’ and ‘cffi’ operate at the C ABI level and have no knowledge of C++ exceptions; any exception that escapes an ‘extern "C"’ boundary produces undefined behaviour (a process crash on most platforms). ‘noexcept’ makes this contract explicit and lets the compiler enforce it. The function body must follow this strict sequence:

1. Immediately zero the ‘error_out’ buffer: ‘if (error_out && error_out_len > 0) error_out[0] = ‘\0’;’
2. Declare ‘ArenaFrame _frame;’ as the very next line, before any other local variable. This RAII guard ensures all ‘R_alloc’ arena memory used by ‘{c_function}’ is freed when ‘{c_function}_c’ returns.
3. Construct a ‘SEXP’ local variable for each input parameter by calling the appropriate helper from Rule 3.
4. Call the original function inside a try block that catches ****all**** exception types. ‘RError’ is the expected error path; ‘std::bad_alloc’ can be thrown by SEXP helper ‘malloc’ calls; ‘std::exception’ covers any other standard exception; the catch-all ‘(...)’ ensures nothing escapes regardless. Every catch arm writes to ‘error_out’ and returns:

```
““cpp
SEXP _result;
try { _result = {c_function}(sexp1, sexp2, ...); }
catch (const RError &_e) {
```



```

        if (error_out && error_out_len > 0) {
            strncpy(error_out, _e.what(), error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input SEXP wrappers */
        return;
    }
    catch (const std::bad_alloc &) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, "out of memory", error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input SEXP wrappers */
        return;
    }
    catch (const std::exception &_e) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, _e.what(), error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input SEXP wrappers */
        return;
    }
    catch (...) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, "unknown C++ exception", error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input SEXP wrappers */
        return;
    }
}
'''

```

5. Unpack the returned ‘_result’ SEXP into caller-supplied output arrays using the helpers from Rule 4. For conditional list components (e.g., ‘csplit’), check the list length and set the corresponding ‘has_’ flag.
6. Free all heap-allocated ‘SEXPREF’ wrapper structs created by the Rule 3 helpers (the underlying data arrays are **not** freed -- the caller owns them). Free the returned ‘_result’ SEXP and any sub-SEXPs it contains.

The Python caller checks ‘error_out’ after every call and raises a Python exception if it is non-empty:

```

'''python
error_buf = ctypes.create_string_buffer(1024)
lib.{c_function}_c(..., error_buf, 1024)
if error_buf.value:
    raise RuntimeError(error_buf.value.decode())
'''

```

This is the only safe error-propagation channel across the C ABI boundary when using ‘ctypes’. If native Python exception propagation is required in the future, the entry-point file should be ported to ‘pybind11’ or ‘Cython’, which understand C++ exceptions and can translate them automatically.

****Rule 6 -- ‘extern "C"’ linkage guard.****

Wrap the ‘extern’ declaration, all helper functions, and the entry point in an ‘#ifdef __cplusplus’ / ‘extern "C"’ / ‘#endif’ block so the file compiles correctly whether the package is built as C or C++.

****Rule 7 -- No runtime R linkage.****

The generated file must introduce zero dependencies on ‘libR.so’. All SEXP types and API functions come exclusively from the fake headers.

Output Format Schema

The generated ‘{c_function}_c.c’ file must follow this exact structure, in order:

```

'''c
/* =====
 * {c_function}_c.c -- Pure C entry point for {c_function}()
 *
 * Original R call form:
 *   {r_call_expression from CSV}
 *
 * Parameter map (SEXP -> plain C):
 *   {param1} ({original SEXP type}) -> {plain C type}  [in]
 *   {param2} ({original SEXP type}) -> {plain C type}  [in]
 *   ...
 *   {return component 1} ({SEXP type}) -> {plain C type}  [out]
 *   {return component 2} ({SEXP type}) -> {plain C type}  [out]
 *   error_out  char *  [out]  -- non-empty on any C++ exception
 *   error_out_len  int  [in]   -- capacity of error_out buffer
 *
 * NOTE: ctypes/cffi have no knowledge of C++ exceptions. All exceptions
 * are caught inside this file and written to error_out. The entry point
 * is declared noexcept to enforce this at the compiler level.
 * Python callers must check error_out after every call.
 *
 * Fake headers required: fake_R.h (and transitive includes)
 * Original C source:      {relative path to C file in c_base_folder}
 * R call sites:          {r_file}:{r_line} [, ...]
 * ===== */

#ifdef __cplusplus
extern "C" {
#endif

#include "fake_R.h"
#include <stdlib.h>
#include <string.h>
#include <stdarg.h>

/* --- Original function declaration (resolved at link time) --- */
extern SEXP {c_function}(...);

/* --- Input SEXP construction helpers (static, no-copy) --- */
static inline SEXP make_int_vec(...) { ... }
static inline SEXP make_real_mat(...) { ... }
/* ... one helper per distinct input conversion pattern ... */

```

```

/* --- Output SEXP extraction helpers (static, copy-out) --- */
static inline void extract_int_vec(...) { ... }
static inline void extract_real_mat(...) { ... }
/* ... one helper per distinct output extraction pattern ... */

/* --- Entry point (noexcept: all exceptions caught internally) --- */
void {c_function}_c(
    /* inputs */
    ...,
    /* outputs */
    ...,
    /* error propagation */
    char *error_out, int error_out_len
) noexcept {
    if (error_out && error_out_len > 0) error_out[0] = '\0';
    ArenaFrame _frame;

    /* wrap inputs */
    ...

    /* call original -- catch ALL exception types; nothing may escape */
    SEXP _result;
    try { _result = {c_function}(...); }
    catch (const RError &_e) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, _e.what(), error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input wrappers */
        return;
    }
    catch (const std::bad_alloc &) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, "out of memory", error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input wrappers */
        return;
    }
    catch (const std::exception &_e) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, _e.what(), error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input wrappers */
        return;
    }
    catch (...) {
        if (error_out && error_out_len > 0) {
            strncpy(error_out, "unknown C++ exception", error_out_len - 1);
            error_out[error_out_len - 1] = '\0';
        }
        /* free input wrappers */
    }
}

```

```

        return;
    }

    /* extract outputs */
    ...

    /* free wrappers */
    ...
}

/* Python caller pattern:
 *   error_buf = ctypes.create_string_buffer(1024)
 *   lib.{c_function}_c(..., error_buf, 1024)
 *   if error_buf.value:
 *       raise RuntimeError(error_buf.value.decode())
 */

#ifdef __cplusplus
} /* extern "C" */
#endif
'''

## Output Saving Instructions
- **File Naming:** Save the generated source file as '{c_function}_c.c' (e.g., 'rpart_c.c', 'pred_rpart_c.c', 'rpartexp2_c.c').
- **Output Directory:** Save the file to the user-specified 'output_folder'. If the user did not explicitly provide one, create and use the default directory 'r2py_rpart/c_entry_points/'.

```

F Phase 5B: Fake-Header Migration

F.1 Skill: /modify-c-folder-with-fake-headers

```

---
name: modify-c-folder-with-fake-headers
description: Batch-modifies all C source and header files in a target folder to replace
    real R API includes with fake_R.h, suppress any unfakeable symbol usages by
    commenting them out, and verifies that both the modified package sources and the
    pre-generated raw entry point files compile cleanly with g++ -x c++.
---

# Modify a C Source Folder to Use Fake R Headers

## Description

This command is designed to run **immediately after** 'generate-r-extern-raw-entry-
points'. Given a 'c_folder' containing the package C source files ('.c' and '.h'),
an 'entry_points_folder' containing the pure-C entry point wrappers generated by '
generate-r-extern-raw-entry-points' (e.g., 'rpart_c.c', 'pred_rpart_c.c'), and the
absolute path to 'fake_r_header' (the master 'fake_R.h' generated by 'generate-r-
extern-fake-headers'), your task is to:

```

1. Recursively discover every '.c' and '.h' file under 'c_folder'.
2. Invoke the '@modify-c-file-with-fake-headers' agent for each file to perform in-place modification.
3. After all files have been processed, compile every modified '.c' file from 'c_folder' ****and**** every '.c' file from 'entry_points_folder' with 'g++ -x c++' to verify correctness, and report a per-file and aggregate result.

The goal is a codebase that compiles cleanly without 'libR.so', relying exclusively on the fake R API headers and the C++ standard library. The entry point files are the primary consumers of the modified package sources -- they must be included in the compilation check, since a successfully modified source file that breaks its entry point wrapper is not a successful outcome.

Execution Steps

Step 1: Discover All C Files

Recursively scan 'c_folder' for every file whose name ends with '.c' or '.h'. Collect their absolute paths into an ordered list -- sorted alphabetically. The ordering has no semantic significance for correctness (all modifications are independent), but alphabetical order produces a deterministic, reproducible discovery log.

Log the discovered file list to the console:

```
'''
Discovered <N> file(s) in <c_folder>:
  [1] <absolute_path_1>
  [2] <absolute_path_2>
  ...
'''
```

Step 2: Process All Files in Parallel

Dispatch the '@modify-c-file-with-fake-headers' agent for ****all files simultaneously****.

Each agent invocation depends only on two inputs that are both fixed before this step begins: the file's own on-disk content and the read-only 'fake_R.h' inventory. No agent reads another file's modified output during its own modification pass, so there are no inter-file ordering constraints. For each file, the agent receives the following sub-steps:

1. ****Invoke the modification agent.**** Call the '@modify-c-file-with-fake-headers' agent, passing:
 - 'c_file': the absolute path to the current file.
 - 'fake_r_header': the absolute path to 'fake_R.h' as provided to this command.
2. ****Collect the per-file report.**** Record the structured summary emitted by the agent (headers replaced, lines commented out, unfakeable symbols found, status).
3. ****Non-blocking error handling.**** If the agent fails, times out, or reports a file-level error, log the following and immediately continue to the next file:


```
'''
ERROR: Failed to process <c_file>. Reason: <error_message>. Proceeding to next file.
'''
Do not abort the overall batch.
```

After all files have been processed, print a cumulative processing summary:

```

'''
Processing complete: <N_ok> file(s) modified successfully, <N_warn> with warnings, <
N_err> failed.
'''

### Step 3: Verify Compilation

After Step 2 has completed for all files, verify that every relevant '.c' file compiles
correctly as C++. The verification covers two sets of files:

- **Set A -- modified package sources:** every '.c' file discovered in Step 1 (the
  files just modified in-place from 'c_folder').
- **Set B -- entry point wrappers:** every '.c' file found directly inside '
  entry_points_folder' (non-recursive; the folder is flat by convention). These files
  were generated by 'generate-r-extern-raw-entry-points' and already contain '#
  include "fake_R.h"' and 'extern SEXP' forward declarations. They are not modified
  by this command, but they must compile against the now-modified package sources to
  confirm the full pipeline is consistent.

'.h' files are excluded from direct compilation in both sets -- they are checked
transitively when their including '.c' files are compiled.

#### 3.1 Determine the Include Paths

Let 'fake_headers_dir' = the **directory** containing 'fake_r_header' (i.e., 'dirname(
fake_r_header)'). This directory must be passed to 'g++' via '-I' so that '#include
"fake_R.h"' and all other fake header names resolve correctly.

#### 3.2 Compile Each '.c' File Individually

For each '.c' file in **Set A**, run:

'''bash
g++ -std=c++17 \
  -x c++ \
  -I"<fake_headers_dir>" \
  -I"<c_folder>" \
  -Wall \
  -Wno-unused-variable \
  -Wno-unused-parameter \
  -fsyntax-only \
  "<c_file>"
'''

For each '.c' file in **Set B**, run:

'''bash
g++ -std=c++17 \
  -x c++ \
  -I"<fake_headers_dir>" \
  -Wall \
  -Wno-unused-variable \
  -Wno-unused-parameter \

```

```

    -fsyntax-only \
    "<entry_point_file>"
'''

```

Entry point files do not need '-I"<c_folder>"' because they reference the original package functions via 'extern' declarations rather than by including the package's internal headers.

Flag explanations:

- '-std=c++17' -- required because 'fake_R.h' uses 'thread_local', 'inline' variables, and other C++17 constructs.
- '-x c++' -- instructs 'g++' to treat the '.c' file as C++ source without renaming it.
- '-I"<fake_headers_dir>"' -- resolves '#include "fake_R.h"' and peer fake headers.
- '-I"<c_folder>"' -- (Set A only) resolves project-internal includes such as '"rpart.h"', '"node.h"', '"rpartproto.h"'.
- '-Wall' -- enables broad diagnostics so genuine issues surface immediately.
- '-Wno-unused-variable' and '-Wno-unused-parameter' -- suppress warnings on variables and parameters that are only retained because they were commented in but not actively used after suppression (an expected side-effect of the comment-out approach).
- '-fsyntax-only' -- performs full type-checking and symbol resolution without producing object files or invoking the linker. This is safe to run on individual translation units that depend on symbols from other translation units.

3.3 Record and Report Results

For each compiled file, record one of the following outcomes:

Outcome	Condition
'PASS'	'g++' exits with status 0 and emits no diagnostics.
'WARN'	'g++' exits with status 0 but emits one or more warning lines.
'FAIL'	'g++' exits with a non-zero status (compilation errors present).

If a file produces 'FAIL', capture the full 'g++' stderr output and include it in the report. Do not abort -- continue compiling the remaining files.

After all compilations complete, print a structured verification report:

```

'''
=====
Compilation Verification Report
=====
Fake headers dir    : <fake_headers_dir>
Source folder      : <c_folder>
Entry points folder: <entry_points_folder>
Files checked      : <N> (<N_a> source + <N_b> entry points)
-----
PASS   : <N_pass> file(s)
WARN   : <N_warn> file(s)
FAIL   : <N_fail> file(s)
-----
Per-file results:

```

```

[PASS] <file_1>
[PASS] <file_2>
[WARN] <file_3>
      warning: <g++ warning text>
[FAIL] <file_4>
      error: <g++ error text, first 10 lines>
...
-----
Overall: <PASS | WARNINGS PRESENT | COMPILATION ERRORS DETECTED>
=====
'''

#### 3.4 Failure Triage Guidance

If any file produces 'FAIL', print the following triage checklist immediately after the
report:

'''
Triage checklist for compilation failures:
1. If the error is "undeclared identifier '<name>'":
   -> '<name>' is used in the C source but not defined in any fake header.
   -> Re-invoke @modify-c-file-with-fake-headers for the failing file with
       an updated fake symbol inventory, or manually add '<name>' to the
       appropriate fake header in <fake_headers_dir>.
2. If the error is "no member named '<field>' in 'SEXPREC'":
   -> The SEXPREC struct in SEXP.h / INTSEXP.h is missing a field used by
       this file. Add the field to the struct definition.
3. If the error is "cannot initialize a variable of type 'T' with an rvalue
   of type 'void *'":
   -> A cast is missing in a fake accessor function. The Category B inline
       in the corresponding fake header needs an explicit cast.
4. If the error is a redefinition conflict:
   -> Two fake headers define the same symbol. Add an #ifndef guard around
       the duplicate definition in the later-included header.
5. If the error relates to a commented-out parameter or missing type in a
   function signature:
   -> The @modify-c-file-with-fake-headers agent violated the parameter rule.
       Re-run the agent for that file; the parameter line must never be
       commented out.
'''

```

F.2 Agent: @modify-c-file-with-fake-headers

```

---
name: modify-c-file-with-fake-headers
description: Modifies a single C source or header file in-place to replace all R API
  header includes with a single fake_R.h include, validates every R external symbol
  used in the file against the fake symbol inventory, and comments out (never removes
  ) any statement that references a symbol absent from the fake headers -- while
  never commenting out function parameter declarations.
---

# Modify a C File to Use Fake R Headers

```


Description

When provided with a 'c_file' (absolute path to a '.c' or '.h' file) and a 'fake_r_header' (absolute path to 'fake_R.h'), your task is to modify the file **in-place** so it compiles correctly with 'g++ -x c++' using the fake R API instead of 'libR.so'. Three categories of change are applied:

1. ****Include replacement**** -- every '#include' line that references a real R API header is commented out and replaced by a single '#include "fake_R.h"' directive.
2. ****Symbol validation**** -- every R external symbol used in the file is cross-referenced against the complete fake symbol inventory derived from 'fake_R.h' and all its transitively included headers.
3. ****Unfakeable-symbol suppression**** -- any statement whose only function is to call or assign a symbol that is absent from the fake headers is commented out (not deleted) with an explanatory note. The one inviolable exception: ****lines that are part of a function parameter declaration must never be commented out****, regardless of which symbols appear on them.

If the file contains no R API includes and no R API symbols, it is left unchanged.

The Non-Negotiable Parameter Rule

A line is a **function parameter line** if it appears inside the parameter list of a function definition -- between the opening '(' that immediately follows the function name and the matching ')' that precedes the opening '{' of the function body. This applies equally to:

- Single-line signatures: 'SEXP foo(SEXP x, SEXP y) {'

- Multi-line signatures:

```
'''c
SEXP foo(SEXP x,
         SEXP y,
         int n) {
'''
```

Function parameter lines **must never be commented out**, even when they reference an unfakeable symbol. Instead, append a trailing '/* NOTE: parameter references unfakeable symbol '{symbol}' -- ensure fake_R.h provides a compatible type */' comment to that line.

This rule exists because commenting out a parameter declaration corrupts the function's ABI and prevents all callers from compiling -- a far worse outcome than leaving an unfakeable type in a parameter position.

Recognised R API Headers

The following include patterns identify real R API headers that must be replaced. Both angle-bracket and quoted forms are recognised:

```
'''
<R.h>           "R.h"
<Rinternals.h>  "Rinternals.h"
<Rdefines.h>    "Rdefines.h"
<Rversion.h>    "Rversion.h"
```

```

<Rmath.h>          "Rmath.h"
<R_ext/Rdynload.h> "R_ext/Rdynload.h"
<R_ext/Error.h>    "R_ext/Error.h"
<R_ext/Memory.h>    "R_ext/Memory.h"
<R_ext/RS.h>        "R_ext/RS.h"
<R_ext/Utils.h>     "R_ext/Utils.h"
<R_ext/Arith.h>     "R_ext/Arith.h"
<R_ext/Boolean.h>   "R_ext/Boolean.h"
<R_ext/Print.h>     "R_ext/Print.h"
'''

```

Additionally, any `#include` whose path matches `R_ext/` or starts with `Rinternals`, `Rdefines`, `Rversion`, `Rmath` is treated as an R API header.

`#include <libintl.h>` is **not** an R API header and must **never** be touched, even when it appears inside an `#ifdef ENABLE-NLS` block alongside real R headers.

Execution Steps

Step 1: Build the Fake Symbol Inventory

Read `'fake_r_header'` (the `'fake_R.h'` file). Collect every `#include "..."` directive it contains and recursively read each referenced header from the same directory. For every header file read, extract the names of all symbols defined: `'typedef's`, `'struct'` tags, `'enum'` constants, `#define` macro names, `'inline'` function names, and `'extern'` variable names. Accumulate all of these into a single flat set: the **fake symbol inventory**.

This inventory is the ground truth for Step 4. A symbol is considered **fakeable** if and only if its name appears in this inventory. A symbol is **unfakeable** if it appears in none of the fake headers despite being an R API symbol (i.e., it would have been provided by one of the real R headers listed above).

Important: The following Category E R Interpreter Items are present in the fake headers as function-pointer stubs. They are therefore **fakeable** for compilation purposes (they compile and link; at runtime they throw `'RError'` if no Python callback has been registered). Do **not** comment out their usage sites:

```

- 'eval', 'Rf_eval'
- 'findVar', 'Rf_findVar'
- 'findVarInFrame', 'Rf_findVarInFrame'
- 'install', 'Rf_install'
- 'R_getVar', 'compat_getVar'

```

Only symbols that are **genuinely absent** from the fake symbol inventory are candidates for suppression in Step 4.

Step 2: Read the Target File

Read the entire contents of `'c_file'` into a line-indexed buffer. Do not modify anything yet. Classify every line as one of the following:

```

- R-include line -- matches one of the R API header patterns from the Description.
- Non-R-include line -- any other #include.
- Preprocessor line -- #ifdef, #endif, #define, etc.

```

- **Function signature line** -- part of a function definition's parameter list (apply the parameter rule from the Description; see also Step 3 for detection logic).
- **Code line** -- any other line inside a function body or at file scope.
- **Comment line** -- already wrapped in `/* */` or `/**/`.

Step 3: Replace R Header Includes

Iterate through the line buffer. For each **R-include line**:

1. **Comment it out in place.** Replace the line with:


```
'''c
/* [FAKE_R] <original_line_content_trimmed> */
'''
```

The original text is preserved inside the comment so the replacement is auditable and reversible.
2. **Mark the insertion point.** After processing all consecutive R-include lines in a block, record the line index immediately following the last commented-out include in that block as the **insertion point** for `'fake_R.h'`. If R-include lines appear in more than one location in the file (e.g., one group near the top and one after `'#ifdef ENABLE_NLS'`), use the position immediately after the **first** group as the insertion point.
3. **Insert the fake_R.h include.** At the recorded insertion point, insert the following new line:


```
'''c
#include "fake_R.h" /* replaces all R API headers above */
'''
```

where `"fake_R.h"` is the **basename only**. The full directory path is supplied to the compiler via the `-I` flag during the verification step in the command; do not embed an absolute path in the source file.
4. **Idempotency guard.** Before inserting, scan the entire file for any existing `'#include "fake_R.h"'` or `'#include <fake_R.h>'` line. If one is already present (from a previous run of this agent), do not insert a second copy.
5. **Files with no R includes.** If the file contains no R-include lines at all, skip this step entirely and proceed to Step 4 without any include modifications.

Step 4: Validate and Suppress Unfakeable Symbols

Walk through every non-comment line in the modified buffer (after Step 3). For each line, tokenise it and check each identifier against the fake symbol inventory built in Step 1. If an identifier is both an R API symbol (i.e., it would have been provided by one of the real R headers) and **absent** from the fake symbol inventory, apply the following rules:

Rule 1 -- Detect the containing syntactic context.

Before deciding whether to comment out the line, determine whether it is part of a function parameter list:

- Parse backwards from the suspicious identifier to find the most recently opened `(''` that has not yet been closed by a matching `')`.

- Then parse backwards further to find whether that '(' is immediately preceded by an identifier followed by optional whitespace -- the signature of a function call or function definition.
- If the '(' is part of a **function definition** (i.e., it is followed eventually by '{' at the same brace depth, with no intervening ';'), the line is a **parameter line**. Apply the parameter rule: append a trailing warning comment and do not comment out the line.
- If the '(' is part of a **function call** or the line is a regular statement in a function body, the line is a **code line** and may be commented out.

Rule 2 -- Comment out single-line statements.

If the line is a code line that forms a complete C statement (ends with ';', or is a 'return' statement, or is a stand-alone expression), comment it out by wrapping the entire original line:

```
/*C
/* [UNFAKEABLE: {symbol}] <original_line_content> */
/*
```

Rule 3 -- Comment out multi-line statements.

If the unfakeable symbol appears on a line that is the start of a multi-line statement (opening expression does not end with ';' and the next line continues the expression), gather all lines belonging to that statement and replace them with a block comment:

```
/*C
/* [UNFAKEABLE: {symbol}] begin
<original_line_1>
<original_line_2>
...
[UNFAKEABLE: {symbol}] end */
/*
```

Rule 4 -- Cascade suppression for dangling references.

If a variable is assigned only in a statement that was suppressed by Rule 2 or Rule 3, and the same variable is subsequently used on another line in the same scope, that subsequent line must also be commented out (with the annotation '/* [UNFAKEABLE: cascaded from {symbol}] ... */'). Apply this rule recursively until no active code line references a variable that has only ever been assigned in suppressed statements.

Rule 5 -- Preprocessor blocks.

If the unfakeable symbol appears inside a '#if'/'#ifdef'/'#elif' block (e.g., '#if R_VERSION < R_Version(4, 5, 0)'), evaluate whether the block as a whole should be suppressed. If the block defines a helper function or macro that itself uses only unfakeable items, comment out the entire block from '#if' to the matching '#endif', replacing it with:

```
/*C
/* [UNFAKEABLE: {symbol}] -- entire preprocessor block suppressed
<original_block_content>
*/
/*
```

```

If the block contains a mix of fakeable and unfakeable items, suppress only the
specific lines that reference unfakeable symbols within the block, following Rules
1-4.

**Rule 6 -- Static file-scope variable declarations.**

A line of the form 'static T var;' where 'T' is an unfakeable type is a variable
declaration at file scope -- not a function parameter. It may be commented out. If
doing so creates dangling references within function bodies, apply Rule 4.

### Step 5: Write the Modified File

Overwrite 'c_file' in-place with the modified line buffer. The file's line count must
be identical to the original: no lines are inserted without offsetting, and no
lines are deleted -- only replaced or commented out. (The single '#include "fake_R.
h"' line inserted in Step 3 is the only net addition; it is injected by replacing
one blank line or by shifting subsequent lines down by exactly one position.)

> Implementation note: the simplest approach is to build the modified file as a list of
strings, where each original line is either kept verbatim, replaced by a comment
variant, or in the case of the insertion point, preceded by the new include line.
The resulting file will have at most one more line than the original (the inserted
'#include "fake_R.h"').

### Step 6: Report Changes

After writing, print a structured summary to the console:

'''
FILE: <c_file>
  R headers replaced      : <count>
  - <original_include_1>
  - <original_include_2>
  ...
  fake_R.h inserted at   : line <N> (after <last_replaced_header>)
  Unfakeable symbols     : <count>
  - <symbol>: <file>:<line_numbers> [commented out | parameter -- warned]
  Lines commented out    : <count>
  Status                  : OK | WARNINGS (see above)
'''

If the file required no changes (no R includes and no unfakeable symbols), print:

'''
FILE: <c_file> -- no changes required
'''

```

G Phase 6: R Package Structural Dependency Analysis

G.1 Skill: /analyze-r-folder-dependencies

```
---
```

```

name: analyze-r-folder-dependencies
description: Analyzes all R scripts in a folder to extract structural dependency
            information file by file, categorizing function calls into language, internal, and
            external dependencies.
---

# Analyze an R Folder for Dependencies

## Description

When provided with a target folder, your task is to recursively scan for all R scripts,
    use the '@analyze-r-file-dependencies' agent on each file individually, and save
    the results as structured JSON files.

## Execution Steps

### Step 1: Scan for R Files

Recursively scan the provided target folder and all of its subdirectories for files
    with '.R' or '.r' extensions. Compile a complete list of their absolute file paths.

### Step 2: Process and Analyze

Iterate through the list of identified R files sequentially. For each file:
1. Invoke the '@analyze-r-file-dependencies' agent against the file and record its
    final JSON output.
2. If the command fails or throws an error for a specific file, log the error to the
    console and immediately continue to the next file in your list. Do not halt the
    entire operation.

### Step 3: Save Output

For every successfully analyzed R file, save the resulting JSON output according to
    these strict rules:
* **Naming Convention:** Use the exact name of the original R script, but replace the
    '.R' or '.r' extension with '.json' (e.g., 'data_cleaning.R' becomes 'data_cleaning.
    json').
* **Output Location:** If the user did not specify the output folder, save the new '.
    json' file in 'structural_analysis/{folder_name}/', where the '{folder_name}' is
    the name of the target folder. The relative path must be the same as its
    corresponding original R file.

```

G.2 Agent: @analyze-r-file-dependencies

```

---
name: analyze-r-file-dependencies
description: Analyzes an R script to extract structural dependency information function
            by function, categorizing calls into language, internal, and external dependencies.
---

# Analyze an R File for Dependencies

```

Description

When provided with an R file (‘.R’), your task is to parse the file, identify every user-defined function, and trace all function calls made within their bodies. You must categorize these dependencies into strict groups and output the result as a structured JSON object.

Execution Steps

Step 1: Extract Function Declarations

Scan the target R file and identify every user-defined function declared within it (e.g., ‘my_func <- function(...) { ... }’). Create a list of these function names.

Step 2: Analyze Function Bodies and Trace Calls

For every function identified in Step 1, thoroughly analyze its execution body. Extract every function call made within that body and categorize it into one of the following three mutually exclusive lists:

* **Language Dependencies:**

* **Definition:** Functions provided by Base R (e.g., ‘c()’, ‘lapply()’, ‘print()’) or functions from installed CRAN/Bioconductor packages (e.g., ‘mutate()’, ‘ggplot()’).

* **Heuristic:** Include any function call that includes a package namespace prefix (e.g., ‘dplyr::select’) or any standard function that is not defined locally.

* **Internal Dependencies:**

* **Definition:** Functions that are defined within the exact same R file, or functions that are explicitly defined in other ‘.R’ files within the same project directory.

* **External Dependencies:**

* **Definition:** Foreign language interfaces (typically C, C++, or Fortran) invoked from within the R code.

* **Heuristic:** Look specifically for calls to ‘.Call()’, ‘.C()’, ‘.Fortran()’, ‘.External()’, or ‘useDynLib()’. Extract the name of the compiled routine being invoked (usually the first string argument passed to these interfaces).

Step 3: Generate JSON Output

Construct a JSON object where each key is the name of a function identified in Step 1. The value for each key must be an object containing exactly three arrays of strings: ‘language_dependencies’, ‘internal_dependencies’, and ‘external_dependencies’. Deduplicate all function names within these arrays.

Output Format Schema

You must output valid JSON strictly adhering to this structure. Do not include markdown code blocks if the system expects raw JSON.

```
““json
{
  "function_name_1": {
    "language_dependencies": ["c", "lapply", "dplyr::filter"],
    "internal_dependencies": ["helper_function_within_file"],
```

```

    "external_dependencies": [".Call(\"c_routine_name\")"]
  },
  "function_name_2": {
    "language_dependencies": ["print"],
    "internal_dependencies": ["function_name_1"],
    "external_dependencies": []
  }
}
'''

```

H Phase 8.1: Language Dependency Extraction

H.1 Skill: /extract-r-folder-language-dependencies

```

---
name: extract-r-folder-language-dependencies
description: Recursively processes a folder of R files, utilizing pre-calculated
  structural analysis JSONs to generate detailed CSV reports of language dependency
  locations.
---

# Extract Language Dependencies from an R Folder

## Description

When provided with a target directory containing R source files and a separate
  directory containing their corresponding structural analysis results (in JSON
  format), your task is to process every R script recursively. You will invoke the ‘
  @extract-r-file-language-dependencies‘ agent on each matched file pair (the ‘.R‘
  file and its corresponding ‘.json‘ analysis), and save the extracted dependency
  data strictly as ‘.csv‘ files in a mirrored output directory.

## Execution Steps

### Step 1: Scan for R Source Files

Recursively scan the provided target directory and all of its subdirectories. Compile a
  complete list of absolute file paths for every file with an ‘.R‘ or ‘.r‘ extension.

### Step 2: Map to JSON and Process

Iterate through the list of identified R files sequentially. For each R file:
1. Locate the JSON Input: Find the corresponding dependency analysis JSON file. It
  will be located in the provided structural analysis directory, sharing the exact
  same relative path as the R file, but with the ‘.R‘/‘.r‘ extension replaced by ‘.
  json‘.
2. Execute File-Level Agent: Invoke the ‘@extract-r-file-language-dependencies‘
  agent, providing it with the R file and its paired JSON file to generate the
  dependency data.
3. Error Handling: If the agent fails, throws an error, or if the required JSON
  file is missing, log a clear error to the console for that specific file and

```



```

        immediately proceed to the next R file in the list. Do not halt the batch execution.

### Step 3: Save Output as CSV

For every successfully analyzed file pair, save the resulting output from the agent
according to these strict rules:
* **Data Format:** The output must be saved strictly as a CSV file.
* **Naming Convention:** Use the exact base name of the original R script, but replace
  the '.R' or '.r' extension with '.csv' (e.g., 'data_cleaning.R' becomes '
  data_cleaning.csv').
* **Output Directory:** Maintain the original directory structure. Save the '.csv' file
  into the user-specified output directory, using the exact same relative path as
  the original R file.
* **Default Output Directory:** If the user did not explicitly specify an output
  directory, create a default directory path named 'language_dependency_analysis/{
  target_folder_name}/' (where '{target_folder_name}' is the base name of the target
  R folder) and mirror the relative paths there.

```

H.2 Agent: @extract-r-file-language-dependencies

```

---
name: extract-r-file-language-dependencies
description: Extracts detailed location and usage data for language dependencies in an
  R file, utilizing pre-calculated structural analysis JSON.
---

# Extract Language Dependencies from an R File

## Description

When provided with an R source file ('.R') and its corresponding structural analysis
JSON (generated by the '@analyze-r-file-dependencies' agent), your task is to
locate every instance where those language dependencies are invoked. You must
extract the exact line numbers and textual calls for each dependency, and output
the results as a well-formatted, sorted CSV.

## Execution Steps

### Step 1: Ingest Inputs

Parse the provided structural analysis JSON to identify the 'language_dependencies'
associated with each user-defined function (the JSON keys).

### Step 2: Locate and Extract Invocations

Scan the raw text of the '.R' file. For every function listed in the JSON, find its
declaration in the R file and scan its body for the specific 'language_dependencies
' listed in its JSON array.

For every individual call to these dependencies, extract the following data points:
* **'language_dependency':** The exact string name of the dependency as listed in the
  JSON (e.g., 'c', 'dplyr::filter').

```

```

* **'function_name'**: The name of the user-defined function (the JSON key) where this
    dependency is currently being executed.
* **'line_number'**: The exact line number in the '.R' file where the invocation of the
    dependency begins.
* **'call_body'**: The exact code snippet of the invocation, including its arguments (e
    .g., 'dplyr::filter(df, column == "value")'). If the call spans multiple lines,
    capture the entire complete statement.

### Step 3: Format and Sort Output

Compile the extracted data points into a CSV format with a header row.

## Output Format Schema

You must output valid CSV data strictly adhering to this structure. Include the header
row. Enclose 'call_body' strings in double quotes to prevent formatting errors from
commas or line breaks within the code.

'''csv
language_dependency,function_name,line_number,call_body
c,my_func_1,14,"c(1, 2, 3)"
c,my_func_2,42,"c(\"A\", \"B\")"
dplyr::filter,my_func_1,15,"dplyr::filter(data, val > 0)"
lapply,my_func_1,22,"lapply(list, function(x) {
  x + 1
})"
print,my_func_2,45,"print(\"Processing complete.\")"
'''

```

I Phase 8.2: Conversion Guide Generation

I.1 Skill: /generate-language-dependency-conversion-guides

```

---
name: generate-language-dependency-conversion-guides
description: Orchestrates the generation of conversion guides for translating language
    dependencies from R to Python based on a provided CSV file.
---

# Generate Language Dependency Conversion Guides

## Description

When provided with a target 'base_folder' and a CSV file matching the following schema:
'''csv
language_dependency,file_path,function_name,line_number,call_body
Re,all.R,bkde,78,"Re(fft(kappa*gcounts, TRUE))"
Re,all.R,bkde2D,156,"Re(fft(rp*sp, inverse = TRUE)/(P1*P2))"
Re,all.R,bkfe,232,"Re(fft(kappam*Gcounts, TRUE))"
abs,all.R,dpill,531,abs(th24Q)
any,all.R,locpoly,610,any(bandwidth <= 0)
as.double,all.R,blkest,265,as.double(x)

```

```

as.double,all.R,blkest,265,as.double(y)
as.double,all.R,blkest,267,as.double(xj)
'''

Your task is to orchestrate the batch processing of this CSV. You must extract rows
corresponding to each unique 'language_dependency' into separate CSV-formatted
strings (each including the header row), invoke the '@generate-language-dependency-
conversion-guide' agent for each subset, and save the resulting guides into a
specified output directory.

## Execution Steps

### Step 1: Identify Unique Dependencies

Parse the input CSV file to identify all unique values in the 'language_dependency'
column.

*Note: The CSV is pre-sorted so identical dependencies are grouped together, allowing
for efficient sequential processing.*

### Step 2: Extract and Process Each Dependency

Iterate through the list of identified unique 'language_dependency' values sequentially
. For each unique dependency:
1. **Extract the CSV Subset:** Isolate all rows matching the current '
language_dependency'. Prepend the standard CSV header row to this subset to create
a valid CSV text string.
2. **Execute Conversion Agent:** Invoke the '@generate-language-dependency-conversion-
guide' agent. Pass the 'base_folder' and the extracted CSV string as inputs to
generate the specific conversion guide.
3. **Error Handling:** If the agent fails or throws an error during generation, log a
clear error message to the console specifying the failed 'language_dependency', and
immediately proceed to the next dependency in your list. Do not halt the overall
batch execution.

### Step 3: Save Output as Markdown Files

For every successfully generated conversion guide, save the agent's markdown output to
the file system according to these strict rules:
* **Naming Convention:** Name the output file exactly '{language_dependency}.md' (e.g.,
'Re.md', 'abs.md').
* **Output Directory:** Save the files to the user-specified output directory. If the
user did not explicitly provide one, create and use a default directory path named
'language_dependency_analysis/conversion_guides/'.

```

I.2 Agent: @generate-language-dependency-conversion-guide

```

---
name: generate-language-dependency-conversion-guide
description: Generates a conversion guide for translating a specific language
dependency from R to Python.
---

# Generate Language Dependency Conversion Guide

```

Description

When provided with a target 'base_folder' and a CSV text input matching the following schema:

```
'''csv
language_dependency,file_path,function_name,line_number,call_body
exp,all.R,locpoly,663,"exp(seq(log(hlow), log(hupp), length.out = Q))"
exp,all.R,sdiag,768,"exp(seq(log(hlow), log(hupp), length.out = Q))"
exp,all.R,sstdiag,845,"exp(seq(log(hlow), log(hupp), length.out = Q))"
'''
```

Your task is to generate a comprehensive conversion guide for the specified 'language_dependency' (e.g., 'exp'). The guide must include:

1. A clear explanation of the 'language_dependency's core functionality in R.
2. A detailed analysis of its usage within the provided CSV data, including the surrounding context retrieved directly from the source R files.
3. A conversion guide to equivalent Python functions/methods, providing general Python code snippets for all scenarios in the CSV.

Execution Steps

Step 1: Ingest Inputs and Gather Context

- Iterate through each row in the CSV. Using the provided 'base_folder', navigate to the exact 'file_path' and 'line_number'.
- Read the target line and the surrounding code blocks. You must read enough lines to fully capture multi-line function calls, identify the data types of arguments being passed (e.g., scalars vs. vectors), and understand how the result interacts with adjacent logic.
- If the R dependency's behavior is complex or ambiguous, search '<https://www.rdocumentation.org/>' to confirm its official functionality, default arguments, and edge cases.
- If there is ambiguity about the valid types of the functions' parameters or return values, you can also look them up in the document '<https://cran.r-project.org/web/packages/rpart/refman/rpart.html>'.

Step 2: Determine Python Equivalents

- ****Prioritize Vectorization:**** Because R inherently vectorizes operations, you must default to evaluating 'numpy', 'scipy', or 'pandas' as the primary Python equivalents.
- Choose the library that best matches R's vectorized nature. For example, if translating R's 'exp', strongly prefer 'numpy.exp()' over the standard library 'math.exp()' unless the context explicitly proves only a scalar is being processed.
- Carefully map R arguments to Python kwargs (e.g., translating R's 'seq(..., length.out=Q)' to Python's 'np.linspace()').
- You are free to search online documentation for the chosen Python libraries to ensure you are using the most efficient and idiomatic functions.

Output Format Schema

You must output the final conversion guide strictly in well-structured Markdown format. Use the following outline to ensure consistency and readability:

```

### 1. Overview of '{language_dependency}' in R

*Provide a concise explanation of what the dependency does, its typical inputs, and
  expected outputs.*

### 2. Contextual Usage Analysis

*Summarize how the dependency is applied across the provided dataset. Mention the data
  types involved and any recurring patterns found in the source files.*

### 3. Python Conversion Strategy

*State the chosen Python library (e.g., 'numpy') and explain why it is the most robust
  equivalent for these specific usage contexts.*

### 4. Step-by-Step Conversion Examples

*For every functional-distinct usage found in the CSV, provide a subsection containing
  :*
- **Locations:** Files and Function names.
- **Original R Context:** The types of input parameters and return values, and a
  generalized code snippet for demonstration.
- **Python Equivalent:** A complete, executable Python code snippet demonstrating the
  converted logic.
- **Explanation:** A brief breakdown of the syntax translation, addressing any nuances
  like zero-based indexing, argument mapping, or library imports.

```

J Phase 9.1: Batch R-to-Python Conversion

J.1 Skill: /convert-r-folder-to-python

```

---
name: convert-r-folder-to-python
description: Converts all functions in all R files within a folder to Python by
  scanning recursively and sequentially invoking /convert-r-file-to-python for each
  file.
---

# Convert All R Files in a Folder to Python

## Description

Translate every R file within a target directory into Python. You will be given the
  path to the R source folder, a folder containing pre-computed JSON dependency maps
  (one per R file, produced by '/analyze-r-folder-dependencies'), a CSV dependency
  graph, and a language dependency conversion guide folder. For each R file
  discovered, you must derive its corresponding JSON map and invoke '/convert-r-file-
  to-python' to handle the per-file translation.

## Parameters

| Parameter | Required | Description |

```

```

|---|---|---|
| 'r_folder' | Yes | Path to the folder containing the R source files to convert. |
| 'json_folder' | Yes | Path to the folder containing the per-file JSON dependency maps
    (e.g., 'structural_analysis/R/'). Each JSON file must share the same base name as
    its corresponding R file (e.g., 'all.R' -> 'all.json'). |
| 'csv_dependency_graph' | Yes | Path to the CSV dependency graph file (e.g., '
    structural_analysis/dependency_levels.csv') that maps functions to their inter-file
    dependency levels. |
| 'language_guides_folder' | Yes | Path to the folder containing language dependency
    conversion guides (one '.md' file per R language construct). |
| 'target_output_folder' | No | Root folder for all converted output. Defaults to '
    conversion_results/R/'. Each R file's functions will be written to '{
    target_output_folder}/{R_file_name}/' as per '/convert-r-file-to-python'
    conventions. |

## Execution Steps

### Step 1: Discover All R Files

- Recursively scan 'r_folder' and all of its subdirectories.
- Identify every file whose name ends with '.R' or '.r'.
- Compile a complete, deterministic list of their absolute paths, sorted alphabetically
.
- Log the discovered file list to the console:

'''
Discovered <N> R file(s) in <r_folder>:
  [1] <absolute_path_1>
  [2] <absolute_path_2>
  ...
'''

### Step 2: Resolve JSON Dependency Maps

For each R file discovered in Step 1:

1. Derive the expected JSON dependency map path by replacing the file's '.R' or '.r'
   extension with '.json' and looking it up within 'json_folder'. Use only the base
   filename -- do not replicate subdirectory structure inside 'json_folder' unless the
   JSON files are organized the same way.
   - *Example:* '{r_folder}/all.R' -> '{json_folder}/all.json'
2. Verify that the derived JSON file exists on disk.
3. If the JSON file does not exist, log a warning and mark that R file as skipped
   **:
   '''
   WARNING: No JSON dependency map found for <r_file> (expected <jjson_path>). Skipping.
   '''
4. Retain only files for which a corresponding JSON map was successfully located. These
   form the execution queue.

### Step 3: Sort the Execution Queue by Dependency Level

Before executing, sort the execution queue to respect inter-file dependency order.

```

1. Parse the CSV dependency graph at 'csv_dependency_graph'.
2. For each R file in the execution queue, match rows by the file's **base name** against the 'r_file' column.
3. For each matching row, extract the numeric level by stripping the optional ' (leaf)' suffix (e.g., '"4 (leaf)"' -> '4').
4. Compute the **maximum level** across all matched rows for each file. If a file has no rows in the CSV, treat its maximum level as '0'.
5. Sort the execution queue in **descending order of maximum level** (files with the highest-level functions -- deepest callees -- are processed first, because other files depend on their outputs). Break ties alphabetically.
6. Log the sorted execution order to the console:

```
'''
Execution order (sorted by max dependency level, highest first):
  [1] <r_file_1>  (max level: <L>)
  [2] <r_file_2>  (max level: <L>)
  ...
'''
```

Step 4: Execute Sequential Conversions

Iterate through the **sorted** execution queue from Step 3 one file at a time. For each R file:

1. **Invoke '/convert-r-file-to-python'**, passing the following explicit context:
 - The absolute path to the current R source file.
 - The absolute path to the corresponding JSON dependency map (resolved in Step 2).
 - The absolute path to the CSV dependency graph ('csv_dependency_graph').
 - The absolute path to the language dependency conversion guides folder ('language_guides_folder').
 - The target output folder ('target_output_folder'). If not provided by the user, use the default 'conversion_results/R/'.
2. **Non-blocking error handling:** If '/convert-r-file-to-python' fails, throws an error, or exits abnormally for a specific file, output a clear error message to the console and immediately proceed to the next file. Do not halt the batch:


```
'''
ERROR: /convert-r-file-to-python failed for <r_file>. Reason: <error_message>.
Proceeding to next file.
'''
```

Step 5: Report Summary

After all files in the execution queue have been attempted, print a final summary:

```
'''
=====
Convert R Folder to Python -- Summary
=====
R source folder      : <r_folder>
JSON dependency maps : <json_folder>
CSV dependency graph : <csv_dependency_graph>
Language guides      : <language_guides_folder>
Target output folder : <target_output_folder>
```

```

-----
Total R files found   : <N_total>
Skipped (no JSON)    : <N_skipped>
Attempted            : <N_attempted>
Succeeded            : <N_ok>
Failed               : <N_err>
=====
'''

If any file failed, list them explicitly so the user knows which files to retry:

'''
Failed files:
- <r_file_1> Reason: <error_1>
- <r_file_2> Reason: <error_2>
'''

```

J.2 Sub-skill: /convert-r-file-to-python

```

---
name: convert-r-file-to-python
description: Converts all functions in an R file to Python by parsing dependency graphs
            and sequentially invoking the function conversion agent.
---

# Convert an R File to Python

## Description

Translate all functions in an entire R file into Python. You will use the R source file
, a comprehensive JSON dependency map, a CSV dependency graph, and a language
dependency conversion guide folder. You must parse the dependencies, establish the
correct conversion order, and invoke the '@convert-r-function-to-python' agent for
each function.

## Execution Steps

### Step 1: Extract Function Definitions and Dependencies

* Parse the provided JSON file. The top-level keys represent all the function names
  within the target R file.
* For each identified function, extract its corresponding R code definition from the R
  source file.
* Isolate the JSON value associated with each function key to serve as its specific
  dependency map.

### Step 2: Sort Functions by Dependency Order (Leaves to Roots)

* Filter the provided dependency CSV file by the 'r_file' column to isolate entries
  relevant only to the current target file.
* Sort the functions identified in Step 1 topologically, from leaves to roots. Ensure
  that child functions always precede their parent functions in the sorted execution
  list.

```



```

* *Note:* If a child function listed in the CSV is not found in the current R file's
  JSON/source, assume it has been converted elsewhere. Ignore it for the purposes of
  sorting the current file.

### Step 3: Execute Sequential Conversions

Iterate through the sorted function list from Step 2 sequentially. For each function:
1. **Invoke Conversion Agent:** Call '@convert-r-function-to-python'. Provide the
  following context:
  * The extracted R code definition.
  * The extracted JSON dependency map.
  * The folder containing language dependency conversion guides.
  * The target output folder for the specific R file. It should be '{
  target_output_folder}/{R_file_name}/'. For example, '{target_output_folder}/all.R/'.
  If the '{target_output_folder}' is not specified, the default should be '
  conversion_results/R/'.
2. **Error Handling:** If the agent fails or throws an error for a specific function,
  log a clear error message to the console. **Do not halt** the overall batch
  execution; proceed to the next function.

```

J.3 Agent: @convert-r-function-to-python

```

---
name: convert-r-function-to-python
description: Converts an R function to its Python equivalent using a provided JSON
  dependency map and conversion guides.
---

# Convert an R Function to Python

## Description

Rewrite a provided R function into Python. You will receive the R function code, a JSON
  file detailing its dependencies, a folder of language dependency conversion guides
  , and a target output folder. You must maintain the exact logic and functionality
  of the original code.

## Execution Steps

### Step 1: Infer Types and Define Prototype

Identify all parameters and return types by referencing the R function's CRAN
  documentation (e.g., the 'rpart' reference manual 'https://cran.r-project.org/web/packages/rpart/refman/rpart.html').
* Create a Python function prototype using complete type annotations.
* **Never omit any types or subtypes inside type annotations.** For example, 'dict[str,
  int]' or 'np.ndarray[Any, np.dtype[np.float64]]' cannot be simplified as 'dict' or
  'np.ndarray' if the subtypes are known. Also, if there are multiple possible types
  for a parameter, include all of them in the annotation (e.g., 'str | None').
* *Example:* 'def example_function(param1: int, param2: np.ndarray[Any, np.dtype[np.
  float64]]) -> np.ndarray[Any, np.dtype[np.float64]]:'

### Step 2: Resolve Dependencies

```

Analyze the provided JSON dependency map. Handle each dependency category as follows:

- * **Language Dependencies:** Find the corresponding '{language_dependency}.md' files (e.g., 'min.md', 'seq.md') in the conversion guides folder for Python translation strategies and examples. You don't have to follow them strictly, but they should guide your translation of R constructs to Python.
- * **Internal Dependencies:** Locate and review the previously converted definitions in the output folder ('{function_name}.json'). Strictly follow their Python logic to ensure correct integration.
- * **External Dependencies:** Call C entry points directly via the 'cffi' '_lib' object already set up in the module. *Example: Map '.Call(C_rpart, ...)' directly to '_lib.rpart_c(...)'. The available C entry points are 'rpart_c', 'pred_rpart_c', 'xpred_c', 'rpartexp2_c', and 'init_rpcallback_c', all loaded from '_rpart_core.so'. Use the module-level helpers '_iptr', '_dptr', '_cptr', '_int32', and '_float64' for array preparation, and always check errors with '_check_error' after each call.* If you are not familiar with R's '.Call' interface, refer to the R Internals manual: '<https://search.r-project.org/R/refmans/base/html/CallExternal.html>'. If you are not familiar with 'cffi' in Python, refer to the documentation: '<https://cffi.readthedocs.io/en/latest/>'. Feel free to look up the project folder 'r2py_rpart' for the implementation details of the C entry points and examples of how to call C entry points from Python.

Step 3: Rewrite Function Body

Translate the R function body into Python syntax.

- * Strictly preserve the original logic and functional intent.
- * Apply the patterns learned from the language conversion guides.
- * Handle R-to-Python structural differences (e.g., 1-based vs. 0-based indexing, data structures, control flow) appropriately.

Output Format

Output strictly valid JSON to '{function_name}.json' in the output folder. ****Do not wrap the output in markdown code blocks.**** The function body may contain multiple lines, and you should include every line as a separate string in the array. Follow the following structure strictly (even if there are no imports, you must include an empty array for imports):

```

'''json
{
  "imports": [
    "import numpy as np"
  ],
  "function_prototype": "def converted_function(param1: int, param2: np.ndarray[Any, np
    .dtype[np.float64]]) -> np.ndarray[Any, np.dtype[np.float64]]:",
  "function_body": [
    "    # Function body goes here",
    "    pass"
  ]
}
'''

```

K Phase 9.2: Package Assembly

K.1 Skill: /combine-python-functions-into-folder

```
---
name: combine-python-functions-into-folder
description: Combines all converted Python functions from the output of /convert-r-
  folder-to-python into one Python file per R source file, then audits __init__.py
  for correctness and verifies the package builds cleanly.
---

# Combine Converted Python Functions into a Folder

## Description

Given a folder that was produced by '/convert-r-folder-to-python' (containing one
  subdirectory per R source file, each holding per-function JSON files), assemble the
  per-function JSON files for every R file into a unified Python module, name each
  output file after its R counterpart ('.R' -> '.py'), and then audit the package's '
  __init__.py' to ensure all public interfaces are correctly exposed. If the package
  cannot be imported or run cleanly, fix any issues before finishing.

## Parameters

| Parameter | Required | Description |
|---|---|---|
| 'conversion_output_folder' | Yes | Path to the folder produced by '/convert-r-folder-
  to-python'. Contains one subdirectory per R file (e.g., 'conversion_results/R/all.R
  /', 'conversion_results/R/build.R/'). |
| 'python_output_folder' | Yes | Path to the target Python package folder where the
  combined '.py' files will be written (e.g., 'r2py_rpart/r2py_rpart/'). |
| 'python_package_folder' | No | Root of the Python package tree used for import
  resolution and '__init__.py' auditing. Defaults to 'python_output_folder'. |

## Execution Steps

### Step 1: Discover All Conversion Subdirectories

- Scan 'conversion_output_folder' for **immediate subdirectories** (do not recurse
  deeper).
- Each subdirectory name is an R source filename (e.g., 'all.R', 'build.R', 'methods.R
  ').
- A subdirectory qualifies if it contains **at least one '.json' file** directly inside
  it.
- Compile a complete, deterministic list of qualifying subdirectories, sorted
  alphabetically by name.
- Log the discovered list to the console:

'''
Discovered <N> conversion subdirectory(ies) in <conversion_output_folder>:
  [1] <subdir_name_1> (<K> JSON file(s))
  [2] <subdir_name_2> (<K> JSON file(s))
  ...
'''
```

```

If no qualifying subdirectories are found, exit immediately with:

'''
ERROR: No conversion subdirectories with JSON files found in <conversion_output_folder
    >. Nothing to combine.
'''

### Step 2: Derive Target Python File Names

For each qualifying subdirectory discovered in Step 1:

1. Take the subdirectory name (e.g., 'all.R').
2. Strip the '.R' or '.r' extension and append '.py' (e.g., 'all.py').
3. The target Python file path is '{python_output_folder}/{derived_name}' (e.g., 'src/
    mypkg/all.py').
4. Log the mapping:

'''
Mapping:
    <subdir_name_1>  ->  <target_python_path_1>
    <subdir_name_2>  ->  <target_python_path_2>
    ...
'''

### Step 3: Execute Sequential Combinations

Iterate through the sorted list from Step 1 one subdirectory at a time. For each
subdirectory:

1. Invoke '/combine-python-functions-into-file', passing the following explicit
    context:
    - The absolute path to the current conversion subdirectory as the source folder.
    - The absolute path to the derived target Python file (from Step 2) as the target
        file.
2. Non-blocking error handling: If '/combine-python-functions-into-file' fails or
    exits abnormally, output a clear error message and proceed to the next subdirectory.
    Do not halt the batch:

'''
ERROR: /combine-python-functions-into-file failed for <subdir_name>. Reason: <
    error_message>. Proceeding to next subdirectory.
'''

### Step 4: Audit '__init__.py'

After all combinations have been attempted, audit the '__init__.py' file located in '
    python_package_folder' (defaulting to 'python_output_folder' if not specified).

Perform the following checks and apply fixes in-place:

1. Remove stale or unnecessary imports: Any symbol imported or re-exported in '
    __init__.py' that no longer exists in any module inside 'python_package_folder'
    must be removed.

```

2. ****Ensure all public interfaces are exposed:**** For every `.py` file written in Step 3 (excluding `__init__.py` itself), verify that each top-level public function (names not prefixed with `_`) is importable via the package. If any are missing from `__init__.py`, add the appropriate import or `__all__` entry.
3. ****No duplicate entries:**** Deduplicate any import lines or `__all__` entries that appear more than once.
4. ****Preserve existing hand-written content:**** Do not remove comments, conditional logic, or imports that are still valid and used.

Log every change made to `__init__.py`:

```
'''
__init__.py audit:
  Removed : <symbol_or_import> (reason: no longer exists in package)
  Added   : <symbol_or_import> (reason: public interface missing from exports)
  ...
  No changes required.  <- (if everything was already correct)
'''
```

Step 5: Verify the Package Builds and Runs Cleanly

After the audit:

1. ****Import check:**** Attempt to import the package from its directory to verify there are no syntax errors or broken imports. Report the result:


```
'''
Import check: PASSED  <- or FAILED: <error>
'''
```
2. ****If the import fails:**** Diagnose the root cause by reading the error traceback carefully. Apply targeted fixes to the affected files (imports, missing symbols, typos, indentation errors). Re-run the import check after each fix. Repeat until the import succeeds or you have exhausted reasonable automated remediation.
3. ****Report remaining issues:**** If automated remediation cannot fully resolve an import failure, describe exactly which file and line need manual attention.

Step 6: Report Summary

Print a final summary:

```
'''
=====
Combine Python Functions into Folder -- Summary
=====
Conversion output folder : <conversion_output_folder>
Python output folder     : <python_output_folder>
Package folder           : <python_package_folder>
-----
Subdirectories found       : <N_total>
Combinations attempted   : <N_attempted>
Succeeded                 : <N_ok>
Failed                   : <N_err>
```

```

__init__.py changes      : <N_changes>
Import check             : PASSED / FAILED
=====
'''

If any combination step failed, list the affected subdirectories:

'''

Failed subdirectories:
- <subdir_1> Reason: <error_1>
- <subdir_2> Reason: <error_2>
'''

```

K.2 Sub-skill: /combine-python-functions-into-file

```

---
name: combine-python-functions-into-file
description: Combines all converted Python functions (stored as JSON) from a specified
            directory into a unified target Python file, handling import deduplication and
            project integration.
---

# Combine Converted Python Functions into a File

## Description

Given a source folder containing JSON-formatted function data and a target Python file,
your task is to extract, format, and combine these functions into a single Python
script. You must resolve duplicate imports, format the code syntactically, and
integrate it safely into the target project structure.

## Execution Steps

### Step 1: Discover and Map JSON Files

* Scan the specified source folder for all files ending with '.json'.
* The filename (excluding the '.json' extension) represents the original R function
  name.

### Step 2: Extract and Format Function Definitions

For each JSON file, extract the 'imports', 'function_prototype', and 'function_body'.
* **Concatenation:** Combine the 'function_prototype' and the 'function_body' array
  into a single multi-line string.
* **Indentation:** Ensure standard Python indentation. The 'function_prototype' should
  have zero indentation, and every line within the 'function_body' must be indented
  by exactly 4 spaces relative to the prototype.

*Example JSON Structure:*
'''json
{
  "imports": [
    "import numpy as np",

```

```

    "from . import _KernSmooth"
],
"function_prototype": "def rlbin(X: np.ndarray[np.float64], Y: np.ndarray[np.float64]
    ], gpoints: np.ndarray[np.float64], truncate: bool = True) -> dict:",
"function_body": [
    "    n = len(X)",
    "    M = len(gpoints)",
    "    return {\"xcounts\": xcnts, \"ycounts\": ycnts}"
]
}
'''

```

Step 3: Deduplicate and Sort Imports

- * Aggregate all 'imports' arrays from every JSON file.
- * Remove any duplicate import statements.
- * Sort the deduplicated imports logically:
 1. Standard library imports (e.g., 'import math').
 2. Third-party imports (e.g., 'import numpy as np').
 3. Local project imports (e.g., 'from .linbin import linbin').
- * Combine the sorted imports and the formatted function definitions into one final multi-line string.

Step 4: Validate and Output the Combined Program

Do not blindly overwrite the target Python file. You must ensure the generated code aligns with the existing project architecture:

1. ****Structural Scan:**** Analyze the target project's build and configuration files (e.g. ., 'pyproject.toml', 'meson.build', '__init__.py') to verify correct module naming conventions and dependency paths.
2. ****Binding Corrections:**** If the structural scan reveals discrepancies in how external C/Fortran bindings are exposed (e.g., '_KernSmooth' is built under a different path than what the JSON imports suggest), automatically fix the import paths in your combined string to match the reality of the build system.
3. ****Safe Write:**** Output the finalized string into the target Python file. If the file already contains code, integrate the new functions safely without breaking existing syntax or overwriting required project-level configurations.

L Phase 9.3: R Interpreter Item Bridge

L.1 Skill: /add-r-interpreter-items

```

---
name: add-r-interpreter-items
description: Writes C helper entry points for Category E R interpreter items, rebuilds
    the shared library, patches __init__.py with cffi callbacks, rewrites rpartcallback.
    py to use cffi exclusively, and resolves all remaining NotImplementedError stubs to
    make the Python package functionally equivalent to the R package.
---

# Add R Interpreter Items

```

```

## Description

After '/combine-python-functions-into-folder', two layers required for full functional
equivalence are missing:

1. C side -- 'make_real_sexp', 'make_int_sexp', 'make_env_sexp', and 'call_install'
   are referenced in 'init_rpcallback.c.c's caller comment but not yet exported from
   '_rpart_core.so'.
2. Python side -- The converted 'rpartcallback.py' mixes 'ctypes' (for callbacks)
   with 'cffi' (for C calls), which is incompatible. All callbacks must be rewritten
   using 'ffi.callback' to match the rest of the package.

This command writes the missing C file, rebuilds the library, patches '__init__.py',
rewrites 'rpartcallback.py', and resolves all 'NotImplementedError' stubs so that
no stub is silently reachable at runtime.

## Parameters

| Parameter | Required | Description |
|---|---|---|
| 'fake_guides_folder' | Yes | Path to the fake header guides (e.g., 'r_extern_analysis
   /fake_guides/'). Provides Python Interop Notes for each Category E item. |
| 'c_entry_points_folder' | Yes | Folder containing existing C entry point files (e.g.,
   'r2py_rpart/c_entry_points/'). The new helper file is written here. |
| 'python_package_folder' | Yes | Python package folder containing '__init__.py' and
   the converted '.py' modules (e.g., 'r2py_rpart/r2py_rpart/'). |
| 'r2py_rpart_root' | Yes | Root of the installable Python package (e.g., 'r2py_rpart
   /'), used for building via 'pip install --no-build-isolation .' |

## Execution Steps

### Step 1: Identify Inputs

1. Category E guides -- In 'fake_guides_folder', find guides whose opening
   blockquote (line 3) begins with '**R Interpreter Item.' (strict prefix match).
   Exclude any guide whose opening blockquote contains 'best-effort fakeable without a
   Python function pointer'. For rpart this yields 'eval.md', 'findVar.md', '
   findVarInFrame.md'. ('install.md' is excluded -- its built-in C++ hash-map requires
   no Python bridge in normal use.)

2. Existing entry points -- Collect the full text of 'init_rpcallback.c.c' from '
   c_entry_points_folder' (contains the caller-comment protocol for 'make_real_sexp'
   etc.) and one representative existing entry point (e.g., 'rpart.c.c') for style
   reference.

3. Existing package files -- Collect the full text of '__init__.py' and '
   rpartcallback.py' from 'python_package_folder'. Also grep every '.py' file in '
   python_package_folder' for 'NotImplementedError' and 'ctypes' to identify all stubs
   and FFI mismatches.

Log a discovery summary:
'''
Category E guides found : eval.md, findVar.md, findVarInFrame.md
Files to patch          : __init__.py, rpartcallback.py

```



```

NotImplementedError hits: <N> across <files>
ctypes usage hits      : <N> across <files>
'''

```

Step 2: Invoke the Agent

Invoke the '@add-r-interpreter-items' agent exactly once, providing all of the following as context:

- The three Category E guide texts (full content of 'eval.md', 'findVar.md', 'findVarInFrame.md')
- The full text of 'init_rpcallback.c.c'
- The full text of one existing entry point file (style reference)
- The full text of 'fake_R.h' (located in 'r2py_rpart/r_fake_headers/' or adjacent)
- The full text of '__init__.py'
- The full text of 'rpartcallback.py' (converted module in 'python_package_folder')
- The grep results for 'NotImplementedError' and 'ctypes' across all '.py' files
- The paths: 'c_entry_points_folder', 'python_package_folder', 'r2py_rpart_root'

The agent handles all file writes, the build step, and all verification. Do not proceed past the agent invocation until it returns.

Step 3: Report Summary

Print the agent's reported results in the following format:

```

'''
=====
Add R Interpreter Items -- Summary
=====
C file written      : c_entry_points/interpreter_helpers_c.c
meson.build updated : YES / NO
Build               : PASSED / FAILED
New symbols verified : make_real_sexp, make_int_sexp,
                      make_env_sexp, call_install
__init__.py patched : YES / NO (<N> additions)
rpartcallback.py fixed : YES / NO (ctypes -> cffi)
NotImplementedError stubs resolved : <N>
NotImplementedError stubs remaining (intentional) : <N>
Import check        : PASSED / FAILED
=====
'''

```

If the build or import check failed, print the full error output and do not mark the step as succeeded. If any 'NotImplementedError' stubs remain, each must be listed with its file, function, and a one-line explanation of why it is intentional (e.g., interactive terminal function with no Python equivalent).

L.2 Agent: @add-r-interpreter-items

```

---
name: add-r-interpreter-items

```

```
description: Writes the C interpreter-item helper entry point file, updates meson.build
, rebuilds _rpart_core.so, patches __init__.py with cffi-based callbacks and ffi.
cdef declarations, rewrites rpartcallback.py to use cffi exclusively, and resolves
all NotImplementedError stubs to make the Python package functionally equivalent to
the R package.
```

```
---
```

```
# Add R Interpreter Items
```

```
## Description
```

```
You receive:
```

- Full text of 'eval.md', 'findVar.md', 'findVarInFrame.md' (Category E fake guides, especially their Section 4 Python Interop Notes)
- Full text of 'init_rpcallback.c.c' (contains the caller-comment protocol for C helpers)
- Full text of one existing entry point file (style reference)
- Full text of 'fake_R.h' (defines SEXPREC struct layout and SEXPTYPE constants)
- Full text of '__init__.py' (cffi setup, existing 'ffi.cdef', existing helpers)
- Full text of 'rpartcallback.py' (currently broken: mixes ctypes callbacks with cffi C calls)
- Grep results for 'NotImplementedError' and 'ctypes' across all '.py' files
- Paths: 'c_entry_points_folder', 'python_package_folder', 'r2py_rpart_root'

```
Your task is to close both the C and Python gaps so the package is fully functional.
```

```
## Execution Steps
```

```
### Step 1: Write 'interpreter_helpers_c.c'
```

```
Read 'fake_R.h' to confirm the exact 'SEXPREC' field order ('type', 'length', 'nrow', 'ncol', 'data'). Read one existing entry point file to match the include order, extern-C bracketing, and noexcept style.
```

```
Write '{c_entry_points_folder}/interpreter_helpers_c.c' exporting four functions. The file must follow the exact same structure as the other entry points: include 'fake_R.h' before the 'extern "C"' block, wrap all exported functions in 'extern "C" { ... }'.
```

```
/**void *make_real_sexp(void *data, int n)**
```

```
Heap-allocates a 'SEXPREC' (via 'malloc') with 'type=REALSXP', 'length=n', 'nrow=n', 'ncol=1', 'data=data'. The data buffer is owned by the caller and is not copied. Returns the SEXPREC pointer cast to 'void *'. On malloc failure, writes to a 'static thread_local char[256]' error buffer and returns 'nullptr'; the buffer is exposed via a companion 'const char *get_make_sexp_error()' function. Does not use 'RError' (this function is called from Python, not from within the C evaluation loop).
```

```
/**void *make_int_sexp(void *data, int n)**
```

```
Identical to 'make_real_sexp' but with 'type=INTSXP'.
```

```
/**void *make_env_sexp(void)**
```

```
Allocates a minimal 'SEXPREC' with 'type=ENVSXP', 'length=0', 'nrow=0', 'ncol=0', 'data=nullptr'. The returned pointer is a stable unique identity key for use as a 'rho'
```

handle in the frame registry. On malloc failure, writes to the same error buffer and returns 'nullptr'.

```
**void *call_install(const char *name)**
```

Calls 'Rf_install(name)' and returns the resulting SYMSXP pointer cast to 'void *'.

This exposes the built-in C++ symbol cache to Python so that Python can pre-populate the frame registry with the exact same pointer keys that 'install()' will produce inside 'init_rpcallback'. Safe to call from Python at any time after library load; does not require 'ArenaFrame'.

Also add 'const char *get_make_sexp_error(void)' that returns the static error buffer, so Python can check for allocation failures.

Step 2: Update 'meson.build'

Open 'meson.build'. Find the 'rpart_src = files(...)' list. Add the line:

```
'''
    'c_entry_points/interpreter_helpers.c',
'''
```

after the existing 'c_entry_points/init_rpcallback.c.c' entry. Do not change any other part of the file.

Step 3: Build and Verify

Run from 'r2py_rpart_root':

```
'''bash
pip install --no-build-isolation .
'''
```

If the build fails, read the compiler error output, fix 'interpreter_helpers.c.c', and retry. Do not proceed until the build passes.

After a successful build, verify all five new symbols are exported:

```
'''bash
nm -D $(python -c "import r2py_rpart, os; print(os.path.dirname(r2py_rpart.__file__))")
/_rpart_core.so \
| grep -E "make_real_sexp|make_int_sexp|make_env_sexp|call_install|
get_make_sexp_error"
'''
```

All five must appear. If any are missing, add '-fkeep-inline-functions' to the compile args for those symbols or restructure the function to have external linkage.

Step 4: Patch '__init__.py' -- 'ffi.cdef' Additions

Append the following block inside the existing 'ffi.cdef("""...""")' call, immediately before the closing '""')'. Do not modify any existing declaration:

```
'''c
/* ---- Interpreter-item helpers (Category E support) ----- */
void *make_real_sexp(void *data, int n);
void *make_int_sexp(void *data, int n);
void *make_env_sexp(void);
void *call_install(const char *name);
const char *get_make_sexp_error(void);
```

```

'''
### Step 5: Patch '__init__.py' -- Module-Level Callback Infrastructure

Insert the following block into '__init__.py' immediately after the '_lib = ffi.dlopen(
    _find_lib())' line. This registers the 'install', 'findVarInFrame', and 'findVar'
callbacks at module load time. The 'eval' callback is intentionally excluded here
-- it is registered per-call inside 'rpartcallback()' because it captures user-
supplied split functions.

'''python
# -----
# Interpreter-item callback infrastructure (Category E, method=4 support)
# -----

# Symbol intern cache: name (str) -> integer address of interned SYMSXP.
# Populated by _py_install; used to pre-populate _frame_registry.
_symbol_handles: dict[str, int] = {}
_symbol_bufs: dict[str, Any] = {}      # keeps ffi buffer objects alive

# Frame registry: rho_ptr (int) -> {sym_ptr (int) -> sexp_ptr (int)}.
# Populated by rpartcallback() before each init_rpcallback_c call.
_frame_registry: dict[int, dict[int, int]] = {}

# Persistent storage for module-level cffi callbacks (prevents GC).
_interp_callbacks: list[Any] = []

# Sentinel value returned by findVarInFrame when a variable is not found.
_R_UNBOUND: int = int(ffi.cast("uintptr_t", _lib.get_R_UnboundValue()))

@ffi.callback("void *(const char *)")
def _install_cb(name_bytes: Any) -> Any:
    name = ffi.string(name_bytes).decode()
    if name not in _symbol_handles:
        buf = ffi.new("char[1]")
        _symbol_bufs[name] = buf
        _symbol_handles[name] = int(ffi.cast("uintptr_t", buf))
    return ffi.cast("void *", _symbol_handles[name])

@ffi.callback("void *(void *, void *)")
def _findVarInFrame_cb(rho: Any, sym: Any) -> Any:
    rho_key = int(ffi.cast("uintptr_t", rho))
    sym_key = int(ffi.cast("uintptr_t", sym))
    val = _frame_registry.get(rho_key, {}).get(sym_key)
    return ffi.cast("void *", val) if val is not None else ffi.cast("void *",
        _R_UNBOUND)

@ffi.callback("void *(void *, void *)")
def _findVar_cb(sym: Any, rho: Any) -> Any:
    # The inherits=TRUE path through compat_getVar is dead code for all
    # standard rpart methods; this stub returns R_UnboundValue safely.

```

```

    return ffi.cast("void *", _R_UNBOUND)

_lib.register_install_fn(_install_cb)
_lib.register_findVarInFrame_fn(_findVarInFrame_cb)
_lib.register_findVar_fn(_findVar_cb)

_interp_callbacks.extend([_install_cb, _findVarInFrame_cb, _findVar_cb])
'''

### Step 6: Rewrite 'rpartcallback.py'

The existing converted 'rpartcallback.py' is broken because it uses 'ctypes.CFUNCTYPE'
for callback creation while the library is loaded via 'cffi'. Rewrite the entire
file so it uses only 'cffi'.

The rewritten function must preserve all input validation and all 'eval1'/'eval2'
computation logic from the existing conversion -- those are correct. Replace only
the C-interface section (Steps 1-9 in the existing comments) with cffi equivalents
following the rules below.

**Replace ctypes constructs with cffi equivalents:**

| Broken pattern (ctypes) | Correct pattern (cffi) |
|---|---|
| 'ctypes.CFUNCTYPE(ret, *args)' | 'ffi.callback("ret (args)")' |
| 'ctypes.cast(x, ctypes.c_void_p).value' | 'int(ffi.cast("uintptr_t", x))' |
| '_SEXPREC(ctypes.Structure)' local struct | '_lib.make_real_sexp(...)' / '_lib.make_int_sexp(...)' |
| '_lib.register_X_fn.restype = None' (ctypes API on cffi lib) | Remove -- cffi does not use '.restype' |
| '_lib.register_X_fn.argtypes = [...]' (ctypes API on cffi lib) | Remove -- cffi resolves from 'ffi.cdef' |
| '_lib.register_X_fn(cb)' | '_lib.register_X_fn(cb)' (valid in cffi too, keep) |

**SEXP construction:** Use the new C helpers instead of local ctypes structs:
'''python
sexp_y = _lib.make_real_sexp(ffi.cast("void *", rho["yback"].ctypes.data), rho["yback"].size)
sexp_w = _lib.make_real_sexp(ffi.cast("void *", rho["wback"].ctypes.data), rho["wback"].size)
sexp_x = _lib.make_real_sexp(ffi.cast("void *", rho["xback"].ctypes.data), rho["xback"].size)
sexp_n = _lib.make_int_sexp(ffi.cast("void *", rho["nback"].ctypes.data), rho["nback"].size)
'''

After each call, check 'ffi.string(_lib.get_make_sexp_error())' and raise 'RuntimeError'
if non-empty.

**Symbol pointers:** Use the module-level '_symbol_handles' dict (already populated by
'_install_cb') rather than a local '_py_install':
'''python
# Trigger install() for each name so _symbol_handles is populated.
for name in (b"yback", b"wback", b"xback", b"nback"):

```

```

    _lib.call_install(name)
'''
Then use '_symbol_handles["yback"]' etc. as the frame registry keys.

**Rho handle:** Use 'make_env_sexp()' instead of a ctypes 1-byte placeholder:
'''python
_rho_sexp = _lib.make_env_sexp()
_rho_ptr = int(ffi.cast("uintptr_t", _rho_sexp))
'''

**Frame registry population:** Write directly to '_frame_registry' (module-level dict
    imported from '__init__.py', or accessed as 'from r2py_rpart import _frame_registry
    '):
'''python
_frame_registry[_rho_ptr] = {
    _symbol_handles["yback"]: int(ffi.cast("uintptr_t", sexp_y)),
    _symbol_handles["wback"]: int(ffi.cast("uintptr_t", sexp_w)),
    _symbol_handles["xback"]: int(ffi.cast("uintptr_t", sexp_x)),
    _symbol_handles["nback"]: int(ffi.cast("uintptr_t", sexp_n)),
}
'''

**Expr handles:** Use 'ffi.new("char[1]")' instead of ctypes 1-byte buffers:
'''python
_expr1_buf = ffi.new("char[1]")
_expr2_buf = ffi.new("char[1]")
_expr1_ptr = int(ffi.cast("uintptr_t", _expr1_buf))
_expr2_ptr = int(ffi.cast("uintptr_t", _expr2_buf))
'''

**Eval callback:** Define using 'ffi.callback'. The callback must catch all Python
    exceptions and return 'ffi.NULL' on failure (cffi will not propagate Python
    exceptions across the C boundary; returning NULL causes the C-side 'isReal(NULL)'
    check to fail with a readable 'RError'):
'''python
_eval_result_keeper: list[Any] = []

@ffi.callback("void *(void *, void *)")
def _eval_cb(expr: Any, rho: Any) -> Any:
    try:
        expr_p = int(ffi.cast("uintptr_t", expr))
        if expr_p == _expr2_ptr:
            result = eval2()
        else:
            result = eval1()
        result = np.ascontiguousarray(result, dtype=np.float64)
        sexp = _lib.make_real_sexp(ffi.cast("void *", result.ctypes.data), result.size)
        _eval_result_keeper.append(result)    # keep numpy array alive
        return sexp
    except Exception:
        return ffi.NULL

_lib.register_eval_fn(_eval_cb)
'''

```

```

**Keep-alive:** Store all cffi objects that must outlive the function call in 'rho["
    _cffi_keep_alive"]':
'''python
rho["_cffi_keep_alive"] = [
    _eval_cb, _expr1_buf, _expr2_buf,
    _rho_sexp, sexp_y, sexp_w, sexp_x, sexp_n,
    _eval_result_keeper,
]
'''
This replaces the '_ctypes_keep_alive' list in the existing conversion.

**Return value:** Keep the same return structure as the existing conversion:
'''python
return {"eval1": eval1, "eval2": eval2, "rho": rho}
'''

### Step 7: Resolve All Remaining 'NotImplementedError' Stubs

Scan every '.py' file in 'python_package_folder' for 'NotImplementedError'. For each
occurrence, classify it and act:

**Class A -- Stubs caused by missing interpreter-item wiring (must be fixed):**
Any stub that says something like "C_init_rpccallback not available", "callbacks not
registered", or similar. These should not exist after Step 6, but verify and fix
any that remain.

**Class B -- 'UseMethod' dispatch stubs (must be fixed):**
Stubs in 'post.py', 'prune.py', and 'meanvar.py' (or wherever 'UseMethod' was converted
to 'NotImplementedError'). Replace each with a direct dispatch based on 'tree.get
("_rpart_class")':
'''python
def post(tree: dict, *args, **kwargs):
    if tree.get("_rpart_class") == "rpart":
        return post_rpart(tree, *args, **kwargs)
    raise TypeError(f"no applicable method for 'post' on object of class "
        f"'{tree.get(\"_rpart_class\", type(tree).__name__)}'")
'''
Apply the same pattern to 'prune' and 'meanvar'.

**Class C -- Interactive or environment-specific stubs (annotate, do not silence):**
Stubs in 'snip_rpart_mouse.py' (uses R's 'identify()' for interactive terminal tree
pruning -- no Python equivalent). Keep 'NotImplementedError' but ensure the message
is specific:
'''python
raise NotImplementedError(
    "snip.rpart.mouse requires interactive terminal graphics (R's identify()); "
    "use snip.rpart() with explicit node indices instead."
)
'''

**Class D -- Formula/model.frame stubs (annotate precisely):**
Stubs where 'eval.parent(model.frame(...))' was replaced with 'NotImplementedError'.
Keep the stub but ensure the message names the specific R function and suggests the

```

```

    Python alternative:
    '''python
raise NotImplementedError(
    "model.frame.rpart requires a formula parser (R's model.frame/eval.parent); "
    "pass a pre-built numpy matrix as 'data' to rpart() instead of a formula string."
)
'''

After processing all stubs, verify: every remaining 'NotImplementedError' has a non-
empty, specific string message. A bare 'raise NotImplementedError' with no message
is not acceptable.

### Step 8: Verify

1. **Import check:**
    '''bash
python -c "import r2py_rpart; print('OK')"
```

'''
 Must print 'OK' with no exception.

2. **Callback registration check:**
 '''bash
python -c "import r2py_rpart; assert len(r2py_rpart._interp_callbacks) == 3,
r2py_rpart._interp_callbacks"
'''
 Must pass (three module-level callbacks registered: install, findVarInFrame, findVar
).

3. **No silent stubs:**
 '''bash
grep -rn "raise NotImplementedError()" {python_package_folder}
'''
 Must return no output. Every 'NotImplementedError' must have a message.

4. **No ctypes usage:**
 '''bash
grep -rn "import ctypes\|ctypes\." {python_package_folder}
'''
 Must return no output. All FFI must go through cffi.

Report pass/fail for each check. If any check fails, diagnose and fix before reporting
completion.